AWS Deep Learning Containers

Developer Guide

AWS Deep Learning Containers: Developer Guide

Table of Contents

What are AWS Deep Learning Containers?

Welcome to the User Guide for the AWS Deep Learning Containers.

AWS Deep Learning Containers (Deep Learning Containers) are a set of Docker images for training and serving models in TensorFlow, TensorFlow 2, PyTorch, and Apache MXNet (Incubating). Deep Learning Containers provide optimized environments with TensorFlow and MXNet, Nvidia CUDA (for GPU instances), and Intel MKL (for CPU instances) libraries and are available in the Amazon Elastic Container Registry (Amazon ECR).

AWS Deep Learning Containers | Amazon Web Services

About this guide

This guide helps you set up and use AWS Deep Learning Containers. This guide also covers setting up Deep Learning Containers with Amazon EC2, Amazon ECS, Amazon EKS, and SageMaker. It covers several use cases that are common for deep learning, for both training and inference. This guide also provides several tutorials for each of the frameworks.

- To run training and inference on Deep Learning Containers for Amazon EC2 using MXNet, PyTorch, TensorFlow, and TensorFlow 2, see Amazon EC2 Tutorials (p. 3)
- To run training and inference on Deep Learning Containers for Amazon ECS using MXNet, PyTorch, and TensorFlow, see Amazon ECS tutorials (p. 13)
- Deep Learning Containers for Amazon EKS offer CPU, GPU, and distributed GPU-based training, as well as CPU and GPU-based inference. To run training and inference on Deep Learning Containers for Amazon EKS using MXNet, PyTorch, and TensorFlow, see Amazon EKS Tutorials (p. 36)
- For an explanation of the Docker-based Deep Learning Containers images, the list of available images, and how to use them, see Deep Learning Containers Images (p. 83)
- For information on security in Deep Learning Containers, see Security in AWS Deep Learning Containers (p. 90)
- For a list of the latest Deep Learning Containers release notes, see Release Notes for Deep Learning Containers (p. 97)

Python 2 Support

The Python open source community has officially ended support for Python 2 on January 1, 2020. The TensorFlow and PyTorch community have announced that the TensorFlow 2.1 and PyTorch 1.4 releases will be the last ones supporting Python 2. Previous releases of the Deep Learning Containers that support Python 2 will continue to be available. However, we will provide updates to the Python 2 Deep Learning Containers only if there are security fixes published by the open source community for those versions. Deep Learning Containers releases with the next versions of the TensorFlow and PyTorch frameworks will not include the Python 2 environments.

Prerequisites

You should be familiar with command line tools and basic Python to successfully run the Deep Learning Containers. Tutorials on how to use each framework are provided by the frameworks themselves. However, this guide shows you how to activate each one and find the appropriate tutorials to get started.

Getting Started With Deep Learning Containers

The following sections describe how to use Deep Learning Containers to run sample code from each of the frameworks on AWS infrastructure. For information on using Deep Learning Containers with SageMaker, see the Use Your Own Algorithms or Models with SageMaker Documentation.

Topics

Amazon EC2 Tutorials

This section shows how to run training and inference on Deep Learning Containers for EC2 using MXNet, PyTorch, TensorFlow, and TensorFlow 2.

Before starting the following tutorials, complete the steps in Amazon EC2 setup (p. 3).

Contents

Amazon EC2 setup

In this section, you learn how to set up AWS Deep Learning Containers with Amazon Elastic Compute Cloud.

Complete the following steps to configure your instance:

- Create an AWS Identity and Access Management user or modify an existing user with the following policies. You can search for them by name in the IAM console's policy tab.
  - AmazonECS_FullAccess Policy
  - AmazonEC2ContainerRegistryFullAccess

  For more information about creating or editing an IAM user, see Adding and Removing IAM Identity Permissions in the IAM user guide.
- Launch an Amazon Elastic Compute Cloud instance (CPU or GPU), preferably a Deep Learning Base AMI. Other AMIs work, but require relevant GPU drivers.
- Connect to your instance by using SSH. For more information about connections, see Troubleshooting Connecting to Your Instance in the *Amazon EC2 user guide.*.
- Ensure your AWS CLI is up to date using the steps in Installing the current AWS CLI Version.

- In your instance, run `aws configure` and provide the credentials of your created user.
- In your instance, run the following command to log in to the Amazon ECR repository where Deep Learning Containers images are hosted.

```
aws ecr get-login-password --region us-east-1 | docker login --username AWS --password-stdin 763104351884.dkr.ecr.us-east-1.amazonaws.com
```

For a complete list of AWS Deep Learning Containers, refer to Deep Learning Containers Images (p. 83).

> **Note**
> MKL users: Read the AWS Deep Learning Containers Intel Math Kernel Library (MKL) Recommendations (p. 85) to get the best training or inference performance.

Next steps

To learn about training and inference on Amazon EC2 with Deep Learning Containers, see Amazon EC2 Tutorials (p. 3).

Training

This section shows how to run training on AWS Deep Learning Containers for Amazon EC2 using Apache MXNet (Incubating), PyTorch, TensorFlow, and TensorFlow 2.

For a complete list of Deep Learning Containers, refer to Deep Learning Containers Images (p. 83).

> **Note**
> MKL users: Read the AWS Deep Learning Containers Intel Math Kernel Library (MKL) Recommendations (p. 85) to get the best training or inference performance.

Contents

TensorFlow training

After you log into your Amazon EC2 instance, you can run TensorFlow and TensorFlow 2 containers with the following commands. You must use `nvidia-docker` for GPU images.

- For CPU-based training, run the following.

```
$ docker run -it <CPU training container>
```

- For GPU-based training, run the following.

```
$ nvidia-docker run -it <GPU training container>
```

The previous command runs the container in interactive mode and provides a shell prompt inside the container. You can then run the following to import TensorFlow.

```
$ python
```

```
>> import tensorflow
```

Press Ctrl+D to return to the bash prompt. Run the following to begin training:

```
git clone https://github.com/fchollet/keras.git
```

```
$ cd keras
```

```
$ python examples/mnist_cnn.py
```

Next steps

To learn inference on Amazon EC2 using TensorFlow with Deep Learning Containers, see TensorFlow Inference (p. 8).

Apache MXNet (Incubating) training

To begin training with Apache MXNet (Incubating) from your Amazon EC2 instance, run the following command to run the container:

- For CPU

```
$ docker run -it <CPU training container>
```

- For GPU

```
$ nvidia-docker run -it <GPU training container>
```

In the terminal of the container, run the following to begin training.

- For CPU

```
$ git clone -b v1.4.x https://github.com/apache/incubator-mxnet.git
python incubator-mxnet/example/image-classification/train_mnist.py
```

- For GPU

```
$ git clone -b v1.4.x https://github.com/apache/incubator-mxnet.git
python incubator-mxnet/example/image-classification/train_mnist.py --gpus 0
```

MXNet training with GluonCV

In the terminal of the container, run the following to begin training using GluonCV. GluonCV v0.6.0 is included in the Deep Learning Containers.

- For CPU

```
$ git clone -b v0.6.0 https://github.com/dmlc/gluon-cv.git
```

```
python gluon-cv/scripts/classification/cifar/train_cifar10.py --model resnet18_v1b
```

- For GPU

```
$ git clone -b v0.6.0 https://github.com/dmlc/gluon-cv.git
python gluon-cv/scripts/classification/cifar/train_cifar10.py --num-gpus 1 --model
 resnet18_v1b
```

Next steps

To learn inference on Amazon EC2 using MXNet with Deep Learning Containers, see Apache MXNet (Incubating) Inference (p. 10).

PyTorch training

To begin training with PyTorch from your Amazon EC2 instance, use the following commands to run the container. You must use **nvidia-docker** for GPU images.

- For CPU

```
$ docker run -it <CPU training container>
```

- For GPU

```
$ nvidia-docker run -it <GPU training container>
```

- If you have docker-ce version 19.03 or later, you can use the --gpus flag with docker:

```
$ docker run -it --gpus <GPU training container>
```

Run the following to begin training.

- For CPU

```
$ git clone https://github.com/pytorch/examples.git
$ python examples/mnist/main.py --no-cuda
```

- For GPU

```
$ git clone https://github.com/pytorch/examples.git
$ python examples/mnist/main.py
```

PyTorch distributed GPU training with NVIDIA Apex

NVIDIA Apex is a PyTorch extension with utilities for mixed precision and distributed training. For more information on the utilities offered with Apex, see the NVIDIA Apex website. Apex is currently supported by Amazon EC2 instances in the following families:

- Amazon EC2 P3 Instances
- Amazon EC2 P2 Instances
- Amazon EC2 G4 Instances
- Amazon EC2 G3 Instances

To begin distributed training using NVIDIA Apex, run the following in the terminal of the GPU training container. This example requires at least two GPUs on your Amazon EC2 instance to run parallel distributed training.

```
$ git clone https://github.com/NVIDIA/apex.git && cd apex
$ python -m torch.distributed.launch --nproc_per_node=2 examples/simple/distributed/
distributed_data_parallel.py
```

Amazon S3 Plugin for PyTorch

Deep Learning Containers include a plugin that enables you to use data from an Amazon S3 bucket for PyTorch training.

1. To begin using the Amazon S3 plugin in Deep Learning Containers, check to make sure that your Amazon EC2 instance has full access to Amazon S3. Create an IAM role that grants Amazon S3 access to an Amazon EC2 instance and attach the role to your instance. You can use the AmazonS3FullAccess or AmazonS3ReadOnlyAccess policies.

2. Set up your `AWS_REGION` environment variable with the region of your choice.

   ```
   export AWS_REGION=us-east-1
   ```

3. Use the following commands to run a container that is compatible with the Amazon S3 plugin. You must use **nvidia-docker** for GPU images.

   - For CPU

     ```
     docker run -it 763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-training:1.8.1-
     cpu-py36-ubuntu18.04-v1.6
     ```

   - For GPU

     ```
     nvidia-docker run -it 763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-
     training:1.8.1-gpu-py36-cu111-ubuntu18.04-v1.7
     ```

4. Run the following to test an example.

   ```
   git clone https://github.com/aws/amazon-s3-plugin-for-pytorch.git
   cd amazon-s3-plugin-for-pytorch/examples
   python s3_cv_iterable_shuffle_example.py
   ```

For more information and additional examples, see the Amazon S3 Plugin for PyTorch repository.

Next steps

To learn inference on Amazon EC2 using PyTorch with Deep Learning Containers, see PyTorch Inference .

Inference

This section shows how to run inference on AWS Deep Learning Containers for Amazon Elastic Compute Cloud using Apache MXNet (Incubating), PyTorch, TensorFlow, and TensorFlow 2. You can also use Elastic Inference to run inference with AWS Deep Learning Containers. For tutorials and more information on Elastic Inference, see Using AWS Deep Learning Containers with Elastic Inference on Amazon EC2.

For a complete list of Deep Learning Containers, refer to Deep Learning Containers Images .

Note

MKL users: read the AWS Deep Learning Containers Intel Math Kernel Library (MKL) Recommendations (p. 85) to get the best training or inference performance.

Contents

TensorFlow Inference

To demonstrate how to use Deep Learning Containers for inference, this example uses a simple *half plus two* model with TensorFlow Serving. We recommend using the Deep Learning Base AMI for TensorFlow. After you log into your instance, run the following:

```
$ git clone -b r1.15 https://github.com/tensorflow/serving.git
$ cd serving
$ git checkout r1.15
```

Use the commands here to start TensorFlow Serving with the Deep Learning Containers for this model. Unlike the Deep Learning Containers for training, model serving starts immediately upon running the container and runs as a background process.

- For CPU instances:

```
$ docker run -p 8500:8500 -p 8501:8501 --name tensorflow-inference --mount
 type=bind,source=$(pwd)/tensorflow_serving/servables/tensorflow/testdata/
saved_model_half_plus_two_cpu,target=/models/saved_model_half_plus_two -e
 MODEL_NAME=saved_model_half_plus_two -d <cpu inference container>
```

For example:

```
$ docker run -p 8500:8500 -p 8501:8501 --name tensorflow-inference --mount
 type=bind,source=$(pwd)/tensorflow_serving/servables/tensorflow/testdata/
saved_model_half_plus_two_cpu,target=/models/saved_model_half_plus_two -e
 MODEL_NAME=saved_model_half_plus_two -d 763104351884.dkr.ecr.us-east-1.amazonaws.com/
tensorflow-inference:1.15.0-cpu-py36-ubuntu18.04
```

- For GPU instances:

```
$ nvidia-docker run -p 8500:8500 -p 8501:8501 --name tensorflow-inference --
mount type=bind,source=$(pwd)/tensorflow_serving/servables/tensorflow/testdata/
saved_model_half_plus_two_gpu,target=/models/saved_model_half_plus_two -e
 MODEL_NAME=saved_model_half_plus_two -d <gpu inference container>
```

For example:

```
$ nvidia-docker run -p 8500:8500 -p 8501:8501 --name tensorflow-inference
 --mount type=bind,source=$(pwd)/tensorflow_serving/servables/tensorflow/
testdata/saved_model_half_plus_two_gpu,target=/models/sad_model_half_plus_two -e
 MODEL_NAME=saved_model_half_plus_two -d 763104351884.dkr.ecr.us-east-1.amazonaws.com/
tensorflow-inference:1.15.0-gpu-py36-cu100-ubuntu18.04
```

- For Inf1 instances:

```
$ docker run -id --name tensorflow-inference -p 8500:8500 --device=/dev/neuron0  --cap-
add IPC_LOCK  --mount type=bind,source={model_path},target=/models/{model_name}  -e
 MODEL_NAME={model_name} <neuron inference container>
```

For example:

```
$ docker run -id --name tensorflow-inference -p 8500:8500 --device=/dev/neuron0  --
cap-add IPC_LOCK  --mount type=bind,source={model_path},target=/models/{model_name}
 -e MODEL_NAME={model_name} 763104351884.dkr.ecr.us-west-2.amazonaws.com/tensorflow-
inference-neuron:1.15.4-neuron-py37-ubuntu18.04-v1.1
```

Next, run inference with the Deep Learning Containers.

```
$ curl -d '{"instances": [1.0, 2.0, 5.0]}' -X POST http://127.0.0.1:8501/v1/models/
saved_model_half_plus_two:predict
```

The output is similar to the following:

```
{
    "predictions": [2.5, 3.0, 4.5
    ]
}
```

> **Note**
> If you want to debug the container's output, you can attach to it using the container name, as in
> the following command:
>
> ```
> $ docker attach <your docker container name>
> ```
>
> In this example you used `tensorflow-inference`.

TensorFlow 2 Inference

To demonstrate how to use Deep Learning Containers for inference, this example uses a simple *half plus two* model with TensorFlow 2 Serving. We recommend using the Deep Learning Base AMI for TensorFlow 2. After you log into your instance run the following.

```
$ git clone -b r2.0 https://github.com/tensorflow/serving.git
$ cd serving
```

Use the commands here to start TensorFlow Serving with the Deep Learning Containers for this model. Unlike the Deep Learning Containers for training, model serving starts immediately upon running the container and runs as a background process.

- For CPU instances:

```
$ docker run -p 8500:8500 -p 8501:8501 --name tensorflow-inference --mount
 type=bind,source=$(pwd)/tensorflow_serving/servables/tensorflow/testdata/
saved_model_half_plus_two_cpu,target=/models/saved_model_half_plus_two -e
 MODEL_NAME=saved_model_half_plus_two -d <cpu inference container>
```

For example:

```
$ docker run -p 8500:8500 -p 8501:8501 --name tensorflow-inference --mount
 type=bind,source=$(pwd)/tensorflow_serving/servables/tensorflow/testdata/
saved_model_half_plus_two_cpu,target=/models/saved_model_half_plus_two -e
 MODEL_NAME=saved_model_half_plus_two -d 763104351884.dkr.ecr.us-east-1.amazonaws.com/
tensorflow-inference:2.0.0-cpu-py36-ubuntu18.04
```

- For GPU instances:

```
$ nvidia-docker run -p 8500:8500 -p 8501:8501 --name tensorflow-inference --
mount type=bind,source=$(pwd)/tensorflow_serving/servables/tensorflow/testdata/
saved_model_half_plus_two_gpu,target=/models/saved_model_half_plus_two -e
 MODEL_NAME=saved_model_half_plus_two -d <gpu inference container>
```

For example:

```
$ nvidia-docker run -p 8500:8500 -p 8501:8501 --name tensorflow-inference
 --mount type=bind,source=$(pwd)/tensorflow_serving/servables/tensorflow/
testdata/saved_model_half_plus_two_gpu,target=/models/sad_model_half_plus_two -e
 MODEL_NAME=saved_model_half_plus_two -d 763104351884.dkr.ecr.us-east-1.amazonaws.com/
tensorflow-inference:2.0.0-gpu-py36-cu100-ubuntu18.04
```

Note
Loading the GPU model server may take some time.

Next, run inference with the Deep Learning Containers.

```
$ curl -d '{"instances": [1.0, 2.0, 5.0]}' -X POST http://127.0.0.1:8501/v1/models/
saved_model_half_plus_two:predict
```

The output is similar to the following.

```
{
    "predictions": [2.5, 3.0, 4.5
    ]
}
```

Note
To debug the container's output, you can use the name to attach to it as shown in the following command:

```
$ docker attach <your docker container name>
```

This example used `tensorflow-inference`.

Apache MXNet (Incubating) Inference

To begin inference with Apache MXNet (Incubating), this example uses a pretrained model from a public S3 bucket.

For CPU instances, run the following command.

```
$ docker run -it --name mms -p 80:8080  -p 8081:8081 <your container image id> \
mxnet-model-server --start --mms-config /home/model-server/config.properties \
--models squeezenet=https://s3.amazonaws.com/model-server/models/squeezenet_v1.1/
squeezenet_v1.1.model
```

For GPU instances, run the following command:

```
$ nvidia-docker run -it --name mms -p 80:8080  -p 8081:8081 <your container image id> \
mxnet-model-server --start --mms-config /home/model-server/config.properties \
--models squeezenet=https://s3.amazonaws.com/model-server/models/squeezenet_v1.1/
squeezenet_v1.1.model
```

The configuration file is included in the container.

With your server started, you can now run inference from a different window by using the following command.

```
$ curl -O https://s3.amazonaws.com/model-server/inputs/kitten.jpg
curl -X POST http://127.0.0.1/predictions/squeezenet -T kitten.jpg
```

After you are done using your container, you can remove it using the following command:

```
$ docker rm -f mms
```

MXNet Inference with GluonCV

To begin inference using GluonCV, this example uses a pretrained model from a public S3 bucket.

For CPU instances, run the following command.

```
$ docker run -it --name mms -p 80:8080  -p 8081:8081 <your container image id> \
mxnet-model-server --start --mms-config /home/model-server/config.properties \
--models gluoncv_yolo3=https://dlc-samples.s3.amazonaws.com/mxnet/gluon/gluoncv_yolo3.mar
```

For GPU instances, run the following command.

```
$ nvidia-docker run -it --name mms -p 80:8080  -p 8081:8081 <your container image id> \
mxnet-model-server --start --mms-config /home/model-server/config.properties \
--models gluoncv_yolo3=https://dlc-samples.s3.amazonaws.com/mxnet/gluon/gluoncv_yolo3.mar
```

The configuration file is included in the container.

With your server started, you can now run inference from a different window by using the following command.

```
$ curl -O https://dlc-samples.s3.amazonaws.com/mxnet/gluon/dog.jpg
curl -X POST http://127.0.0.1/predictions/gluoncv_yolo3/predict -T dog.jpg
```

Your output should look like the following:

```
{
  "bicycle": [
    "[ 79.674225  87.403786 409.43515   323.12167 ]",
    "[ 98.69891  107.480446 200.0086    155.13412 ]"
  ],
  "car": [
    "[336.61322    56.533463 499.30566  125.0233   ]"
  ],
  "dog": [
    "[100.50538 156.50375 223.014    384.60873]"
  ]
```

```
}
```

After you are done using your container, you can remove it using this command.

```
$ docker rm -f mms
```

PyTorch Inference

Deep Learning Containers with PyTorch version 1.6 and later use TorchServe for inference calls. Deep Learning Containers with PyTorch version 1.5 and earlier use `mxnet-model-server` for inference calls.

PyTorch 1.6 and later

To run inference with PyTorch, this example uses a model pretrained on Imagenet from a public S3 bucket. Inference is served using TorchServe. For more information, see this blog on Deploying PyTorch inference with TorchServe.

For CPU instances:

```
$ docker run -itd --name torchserve -p 80:8080  -p 8081:8081 <your container image id> \
torchserve --start --ts-config /home/model-server/config.properties \
--models pytorch-densenet=https://torchserve.s3.amazonaws.com/mar_files/densenet161.mar
```

For GPU instances

```
$ nvidia-docker run -itd --name torchserve -p 80:8080  -p 8081:8081 <your container image
 id> \
torchserve --start --ts-config /home/model-server/config.properties \
--models pytorch-densenet=https://torchserve.s3.amazonaws.com/mar_files/densenet161.mar
```

If you have docker-ce version 19.03 or later, you can use the `--gpus` flag when you start Docker.

The configuration file is included in the container.

With your server started, you can now run inference from a different window by using the following.

```
$ curl -O https://s3.amazonaws.com/model-server/inputs/flower.jpg
curl -X POST http://127.0.0.1:80/predictions/pytorch-densenet -T flower.jpg
```

After you are done using your container, you can remove it using the following.

```
$ docker rm -f torchserve
```

PyTorch 1.5 and earlier

To run inference with PyTorch, this example uses a model pretrained on Imagenet from a public S3 bucket. Similar to MXNet containers, inference is served using mxnet-model-server, which can support any framework as the backend. For more information, see Model Server for Apache MXNet and this blog on Deploying PyTorch inference with MXNet Model Server.

For CPU instances:

```
$ docker run -itd --name mms -p 80:8080  -p 8081:8081 <your container image id> \
mxnet-model-server --start --mms-config /home/model-server/config.properties \
```

```
--models densenet=https://dlc-samples.s3.amazonaws.com/pytorch/multi-model-server/densenet/
densenet.mar
```

For GPU instances

```
$ nvidia-docker run -itd --name mms -p 80:8080  -p 8081:8081 <your container image id> \
mxnet-model-server --start --mms-config /home/model-server/config.properties \
--models densenet=https://dlc-samples.s3.amazonaws.com/pytorch/multi-model-server/densenet/
densenet.mar
```

If you have docker-ce version 19.03 or later, you can use the `--gpus` flag when you start Docker.

The configuration file is included in the container.

With your server started, you can now run inference from a different window by using the following.

```
$ curl -O https://s3.amazonaws.com/model-server/inputs/flower.jpg
curl -X POST http://127.0.0.1/predictions/densenet -T flower.jpg
```

After you are done using your container, you can remove it using the following.

```
$ docker rm -f mms
```

Next steps

To learn about using custom entrypoints with Deep Learning Containers on Amazon ECS, see Custom entrypoints (p. 36).

Custom Entrypoints

For some images, Deep Learning Containers uses a custom entrypoint script. If you want to use your own entrypoint, you can override the entrypoint as follows.

- To specify a custom entrypoint script to run, use this command.

```
docker run --entrypoint=/path/to/custom_entrypoint_script -it <image> /bin/bash
```

- To set the entrypoint to be empty, use this command.

```
docker run --entrypoint="" <image> /bin/bash
```

Amazon ECS tutorials

This section shows how to run training and inference on AWS Deep Learning Containers for Amazon ECS using MXNet, PyTorch, and TensorFlow.

Before starting the following tutorials, complete the steps in Amazon ECS setup (p. 14).

For a complete list of Deep Learning Containers, refer to Deep Learning Containers Images (p. 83).

> **Note**
> MKL users: Read the AWS Deep Learning Containers Intel Math Kernel Library (MKL) Recommendations (p. 85) to get the best training or inference performance.

Contents

Amazon ECS setup

This topic shows how to setup AWS Deep Learning Containers with Amazon Elastic Container Service.

Contents

Prerequisites

This setup guide assumes that you have completed the following prerequisites:

- Install and configure the latest version of the AWS CLI. For more information about installing or upgrading the AWS CLI, see Installing the AWS Command Line Interface.
- Complete the steps in Setting Up with Amazon ECS.
- One of the following is true:
  - Your user has administrator access. For more information, see Setting Up with Amazon ECS.
  - Your user has the IAM permissions to create a service role. For more information, see Creating a Role to Delegate Permissions to an AWS Service.
  - A user with administrator access has manually created these IAM roles so that they're available on the account to be used. For more information, see Amazon ECS Service Scheduler IAM Role and Amazon ECS Container Instance IAM Role in the *Amazon Elastic Container Service Developer Guide*.
- The Amazon CloudWatch Logs IAM policy is added to the Amazon ECS Container Instance IAM role, which allows Amazon ECS to send logs to Amazon CloudWatch. For more information, see CloudWatch Logs IAM Policy in the *Amazon Elastic Container Service Developer Guide*.
- Generate a key pair. For more information see Amazon EC2 Key Pairs.
- Create a new security group or update an existing security group to have the ports open for your desired inference server.
  - For MXNet inference, ports 80 and 8081 open to TCP traffic.
  - For TensorFlow inference, ports 8501 and 8500 open to TCP traffic.

  For more information see Amazon EC2 Security Groups.

Setting up Amazon ECS for Deep Learning Containers

This section explains how to set up Amazon ECS to use Deep Learning Containers.

> **Important**
> If your account has already created the Amazon ECS service-linked role, then that role is used by default for your service unless you specify a role here. The service-linked role is required if your task definition uses the **awsvpc** network mode or if the service is configured to use any of the following: Service discovery, an external deployment controller, multiple target groups, or Elastic Inference accelerators. If this is the case, you should not specify a role here. For more information, see Using Service-Linked Roles for Amazon ECS in the *Amazon ECS Developer Guide*.

Run the following actions from your host.

1. Create an Amazon ECS cluster in the Region that contains the key pair and security group that you created previously.

```
aws ecs create-cluster --cluster-name ecs-ec2-training-inference --region us-east-1
```

2. Launch one or more Amazon EC2 instances into your cluster. For GPU-based work, refer to Working with GPUs on Amazon ECS in the *Amazon ECS Developer Guide* to guide your instance type selection. After you select your instance type, select an ECS-optimized AMI that fits your use case. For CPU-based work, you can use the Amazon Linux or Amazon Linux 2 ECS-optimized AMIs. For GPU-based work, you must use the ECS GPU-optimized AMI and a p2/p3 instance type. You can find the Amazon ECS-optimized AMI IDs at Amazon ECS-optimized AMIs. In this example, you launch one instance with a GPU-based AMI with 100 GB of disk size in us-east-1.

   a. Create a file named `my_script.txt` with the following contents. Reference the same cluster name that you created in the previous step.

   ```
   #!/bin/bash
   echo ECS_CLUSTER=ecs-ec2-training-inference >> /etc/ecs/ecs.config
   ```

   b. (Optional) Create a file named `my_mapping.txt` with the following content, which changes the size of the root volume after the instance is created.

   ```
   [
       {
           "DeviceName": "/dev/xvda",
           "Ebs": {
               "VolumeSize": 100
           }
       }
   ]
   ```

   c. Launch an Amazon EC2 instance with the Amazon ECS-optimized AMI and attach it to the cluster. Use the security group ID and key pair name that you created and replace them in the following command. To get the latest Amazon ECS-optimized AMI ID, see Amazon ECS-optimized AMIs in the *Amazon Elastic Container Service Developer Guide*.

   ```
   aws ec2 run-instances --image-id ami-0dfdeb4b6d47a87a2 \
                          --count 1 \
                          --instance-type p2.8xlarge \
                          --key-name key-pair-1234 \
                          --security-group-ids sg-abcd1234 \
                          --iam-instance-profile Name="ecsInstanceRole" \
                          --user-data file://my_script.txt \
                          --block-device-mapping file://my_mapping.txt \
                          --region us-east-1
   ```

   In the Amazon EC2 console, you can verify that this step was successful by the `instance-id` from the response.

You now have an Amazon ECS cluster with container instances running. Verify that the Amazon EC2 instances are registered with the cluster with the following steps.

To verify that the Amazon EC2 instance is registered with the cluster

1. Open the Amazon ECS console at https://console.aws.amazon.com/ecs/.
2. Select the cluster with your registered Amazon EC2 instances.
3. On the **Cluster** page, choose **ECS Instances**.

4.   Verify that the **Agent Connected** value is **True** for the `instance-id` created in previous step. Also, note the CPU available and memory available from the console as these values can be useful in the following tutorials. It might take a few minutes to appear in the console.

Next steps

To learn about training and inference with Deep Learning Containers on Amazon ECS, see Amazon ECS tutorials (p. 13).

Training

This section shows how to run training on AWS Deep Learning Containers for Amazon Elastic Container Service using Apache MXNet (Incubating), PyTorch, TensorFlow, and TensorFlow 2.

For a complete list of Deep Learning Containers, refer to Deep Learning Containers Images (p. 83).

> **Note**
> MKL users: Read the AWS Deep Learning Containers Intel Math Kernel Library (MKL) Recommendations (p. 85) to get the best training or inference performance.

> **Important**
> If your account has already created the Amazon ECS service-linked role, that role is used by default for your service unless you specify a role here. The service-linked role is required if your task definition uses the **awsvpc** network mode or if the service is configured to use service discovery. The role is also required if the service uses an external deployment controller, multiple target groups, or Elastic Inference accelerators in which case you should not specify a role here. For more information, see Using Service-Linked Roles for Amazon ECS in the *Amazon ECS Developer Guide*.

Contents

- TensorFlow training (p. 16)
- Apache MXNet (Incubating) training (p. 18)
- PyTorch training (p. 20)
- Amazon S3 Plugin for PyTorch (p. 22)
- Next steps (p. 25)

TensorFlow training

Before you can run a task on your ECS cluster, you must register a task definition. Task definitions are lists of containers grouped together. The following example uses a sample Docker image that adds training scripts to Deep Learning Containers. You can use this script with either TensorFlow or TensorFlow 2. To use it with TensorFlow 2, change the Docker image to a TensorFlow 2 image.

1.   Create a file named `ecs-deep-learning-container-training-taskdef.json` with the following contents.

- For CPU

```
{
  "requiresCompatibilities": [
    "EC2"
  ],
  "containerDefinitions": [{
    "command": [
      "mkdir -p /test && cd /test && git clone https://github.com/fchollet/keras.git &&
  chmod +x -R /test/ && python keras/examples/mnist_cnn.py"
    ],
```

```
  "entryPoint": [
   "sh",
   "-c"
  ],
  "name": "tensorflow-training-container",
  "image": "763104351884.dkr.ecr.us-east-1.amazonaws.com/tensorflow-inference:1.15.2-
cpu-py36-ubuntu18.04",
  "memory": 4000,
  "cpu": 256,
  "essential": true,
  "portMappings": [{
   "containerPort": 80,
   "protocol": "tcp"
  }],
  "logConfiguration": {
   "logDriver": "awslogs",
   "options": {
    "awslogs-group": "awslogs-tf-ecs",
    "awslogs-region": "us-east-1",
    "awslogs-stream-prefix": "tf",
    "awslogs-create-group": "true"
   }
  }
 }],
 "volumes": [],
 "networkMode": "bridge",
 "placementConstraints": [],
 "family": "TensorFlow"
}
```

- For GPU

```
{
     "requiresCompatibilities": [
       "EC2"
     ],
     "containerDefinitions": [
       {
    "command": [
             "mkdir -p /test && cd /test && git clone https://github.com/fchollet/
keras.git && chmod +x -R /test/ && python keras/examples/mnist_cnn.py"
         ],
         "entryPoint": [
             "sh",
             "-c"
         ],
       "name": "tensorflow-training-container",
       "image": "763104351884.dkr.ecr.us-east-1.amazonaws.com/tensorflow-
training:1.15.2-gpu-py37-cu100-ubuntu18.04",
         "memory": 6111,
         "cpu": 256,
         "resourceRequirements" : [{
           "type" : "GPU",
           "value" : "1"
         }],
         "essential": true,
         "portMappings": [
           {
              "containerPort": 80,
              "protocol": "tcp"
           }
         ],
         "logConfiguration": {
             "logDriver": "awslogs",
             "options": {
```

```
              "awslogs-group": "awslogs-tf-ecs",
              "awslogs-region": "us-east-1",
              "awslogs-stream-prefix": "tf",
          "awslogs-create-group": "true"
              }
          }
        }
    ],
    "volumes": [],
    "networkMode": "bridge",
    "placementConstraints": [],
    "family": "tensorflow-training"
  }
```

2. Register the task definition. Note the revision number in the output and use it in the next step.

```
aws ecs register-task-definition --cli-input-json file://ecs-deep-learning-container-
training-taskdef.json
```

3. Create a task using the task definition. You need the revision number from the previous step and the name of the cluster you created during setup

```
aws ecs run-task --cluster ecs-ec2-training-inference --task-definition tf:1
```

4. Open the Amazon ECS console at https://console.aws.amazon.com/ecs/.
5. Select the `ecs-ec2-training-inference` cluster.
6. On the **Cluster** page, choose **Tasks**.
7. After your task is in a `RUNNING` state, choose the task identifier.
8. Under **Containers**, expand the container details.
9. Under **Log Configuration**, choose **View logs in CloudWatch**. This takes you to the CloudWatch console to view the training progress logs.

Next steps

To learn inference on Amazon ECS using TensorFlow with Deep Learning Containers, see TensorFlow inference (p. 25).

Apache MXNet (Incubating) training

Before you can run a task on your Amazon Elastic Container Service cluster, you must register a task definition. Task definitions are lists of containers grouped together. The following example uses a sample Docker image that adds training scripts to Deep Learning Containers.

1. Create a file named `ecs-deep-learning-container-training-taskdef.json` with the following contents.

   - For CPU

```
{
   "requiresCompatibilities":[
      "EC2"
   ],
   "containerDefinitions":[
      {
         "command":[
            "git clone -b 1.4 https://github.com/apache/incubator-mxnet.git &&
python /incubator-mxnet/example/image-classification/train_mnist.py"
         ],
```

```json
        "entryPoint":[
            "sh",
            "-c"
        ],
        "name":"mxnet-training",
        "image":"763104351884.dkr.ecr.us-east-1.amazonaws.com/mxnet-training:1.6.0-
cpu-py36-ubuntu16.04",
        "memory":4000,
        "cpu":256,
        "essential":true,
        "portMappings":[
            {
                "containerPort":80,
                "protocol":"tcp"
            }
        ],
        "logConfiguration":{
            "logDriver":"awslogs",
            "options":{
                "awslogs-group":"/ecs/mxnet-training-cpu",
                "awslogs-region":"us-east-1",
                "awslogs-stream-prefix":"mnist",
                "awslogs-create-group":"true"
            }
        }
    }
    ],
    "volumes":[

    ],
    "networkMode":"bridge",
    "placementConstraints":[

    ],
    "family":"mxnet"
}
```

- For GPU

```json
{
    "requiresCompatibilities":[
        "EC2"
    ],
    "containerDefinitions":[
        {
            "command":[
                "git clone -b 1.4 https://github.com/apache/incubator-mxnet.git &&
 python /incubator-mxnet/example/image-classification/train_mnist.py --gpus 0"
            ],
            "entryPoint":[
                "sh",
                "-c"
            ],
            "name":"mxnet-training",
            "image":"763104351884.dkr.ecr.us-east-1.amazonaws.com/mxnet-training:1.6.0-
gpu-py36-cu101-ubuntu16.04",
            "memory":4000,
            "cpu":256,
            "resourceRequirements":[
                {
                    "type":"GPU",
                    "value":"1"
                }
            ],
            "essential":true,
```

```
            "portMappings":[
                {
                    "containerPort":80,
                    "protocol":"tcp"
                }
            ],
            "logConfiguration":{
                "logDriver":"awslogs",
                "options":{
                    "awslogs-group":"/ecs/mxnet-training-gpu",
                    "awslogs-region":"us-east-1",
                    "awslogs-stream-prefix":"mnist",
                    "awslogs-create-group":"true"
                }
            }
        }
    ],
    "volumes":[

    ],
    "networkMode":"bridge",
    "placementConstraints":[

    ],
    "family":"mxnet-training"
}
```

2. Register the task definition. Note the revision number in the output and use it in the next step.

    ```
    aws ecs register-task-definition --cli-input-json file://ecs-deep-learning-container-
    training-taskdef.json
    ```

3. Create a task using the task definition. You need the revision number from the previous step.

    ```
    aws ecs run-task --cluster ecs-ec2-training-inference --task-definition mx:1
    ```

4. Open the Amazon ECS console at https://console.aws.amazon.com/ecs/.
5. Select the ecs-ec2-training-inference cluster.
6. On the **Cluster** page, choose **Tasks**.
7. After your task is in a RUNNING state, choose the task identifier.
8. Under **Containers**, expand the container details.
9. Under **Log Configuration**, choose **View logs in CloudWatch**. This takes you to the CloudWatch console to view the training progress logs.

Next steps

To learn inference on Amazon ECS using MXNet with Deep Learning Containers, see Apache MXNet (Incubating) inference (p. 29).

PyTorch training

Before you can run a task on your Amazon ECS cluster, you must register a task definition. Task definitions are lists of containers grouped together. The following example uses a sample Docker image that adds training scripts to Deep Learning Containers.

1. Create a file named ecs-deep-learning-container-training-taskdef.json with the following contents.

    - For CPU

```
{
    "requiresCompatibilities":[
        "EC2"
    ],
    "containerDefinitions":[
        {
            "command":[
                "git clone https://github.com/pytorch/examples.git && python examples/
mnist/main.py --no-cuda"
            ],
            "entryPoint":[
                "sh",
                "-c"
            ],
            "name":"pytorch-training-container",
            "image":"763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-
training:1.5.1-cpu-py36-ubuntu16.04",
            "memory":4000,
            "cpu":256,
            "essential":true,
            "portMappings":[
                {
                    "containerPort":80,
                    "protocol":"tcp"
                }
            ],
            "logConfiguration":{
                "logDriver":"awslogs",
                "options":{
                    "awslogs-group":"/ecs/pytorch-training-cpu",
                    "awslogs-region":"us-east-1",
                    "awslogs-stream-prefix":"mnist",
                    "awslogs-create-group":"true"
                }
            }
        }
    ],
    "volumes":[

    ],
    "networkMode":"bridge",
    "placementConstraints":[

    ],
    "family":"pytorch"
}
```

- For GPU

```
{
    "requiresCompatibilities": [
        "EC2"
    ],
    "containerDefinitions": [
        {
    "command": [
                "git clone https://github.com/pytorch/examples.git && python
 examples/mnist/main.py"
            ],
            "entryPoint": [
                "sh",
                "-c"
            ],
```

```
          "name": "pytorch-training-container",
          "image": "763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-
training:1.5.1-gpu-py36-cu101-ubuntu16.04",
          "memory": 6111,
          "cpu": 256,
          "resourceRequirements" : [{
            "type" : "GPU",
            "value" : "1"
          }],
          "essential": true,
          "portMappings": [
            {
              "containerPort": 80,
              "protocol": "tcp"
            }
          ],
          "logConfiguration": {
              "logDriver": "awslogs",
              "options": {
                  "awslogs-group": "/ecs/pytorch-training-gpu",
                  "awslogs-region": "us-east-1",
                  "awslogs-stream-prefix": "mnist",
                      "awslogs-create-group": "true"
              }
          }
        }
      ],
      "volumes": [],
      "networkMode": "bridge",
      "placementConstraints": [],
      "family": "pytorch-training"
    }
```

2. Register the task definition. Note the revision number in the output and use it in the next step.

    ```
    aws ecs register-task-definition --cli-input-json file://ecs-deep-learning-container-
    training-taskdef.json
    ```

3. Create a task using the task definition. You need the revision identifier from the previous step.

    ```
    aws ecs run-task --cluster ecs-ec2-training-inference --task-definition pytorch:1
    ```

4. Open the Amazon ECS console at https://console.aws.amazon.com/ecs/.

5. Select the `ecs-ec2-training-inference` cluster.

6. On the **Cluster** page, choose **Tasks**.

7. After your task is in a RUNNING state, choose the task identifier.

8. Under **Containers**, expand the container details.

9. Under **Log Configuration**, choose **View logs in CloudWatch**. This takes you to the CloudWatch console to view the training progress logs.

Amazon S3 Plugin for PyTorch

Deep Learning Containers include a plugin that enables you to use data from an Amazon S3 bucket for PyTorch training.

1. To begin using the Amazon S3 plugin in Amazon ECS, set up your AWS_REGION environment variable with the region of your choice.

```
export AWS_REGION=us-east-1
```

2. Create a file named `ecs-deep-learning-container-pytorch-s3-plugin-taskdef.json` with the following contents.

- For CPU

```
{
    "requiresCompatibilities":[
        "EC2"
    ],
    "containerDefinitions":[
        {
            "command":[
                "git clone https://github.com/aws/amazon-s3-plugin-for-pytorch.git &&
 python amazon-s3-plugin-for-pytorch/examples/s3_imagenet_example.py"
            ],
            "entryPoint":[
                "sh",
                "-c"
            ],
            "name":"pytorch-s3-plugin-container",
            "image":"763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-
training:1.8.1-cpu-py36-ubuntu18.04-v1.6",
            "memory":4000,
            "cpu":256,
            "essential":true,
            "portMappings":[
                {
                    "containerPort":80,
                    "protocol":"tcp"
                }
            ],
            "logConfiguration":{
                "logDriver":"awslogs",
                "options":{
                    "awslogs-group":"/ecs/pytorch-s3-plugin-cpu",
                    "awslogs-region":"us-east-1",
                    "awslogs-stream-prefix":"imagenet",
                    "awslogs-create-group":"true"
                }
            }
        }
    ],
    "volumes":[

    ],
    "networkMode":"bridge",
    "placementConstraints":[

    ],
    "family":"pytorch-s3-plugin"
}
```

- For GPU

```
{
    "requiresCompatibilities": [
      "EC2"
    ],
    "containerDefinitions": [
      {
```

```
    "command": [
            "git clone https://github.com/aws/amazon-s3-plugin-for-pytorch.git &&
 python amazon-s3-plugin-for-pytorch/examples/s3_imagenet_example.py"
            ],
            "entryPoint": [
                "sh",
                "-c"
            ],
        "name": "pytorch-s3-plugin-container",
        "image": "763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-
training:1.8.1-gpu-py36-cu111-ubuntu18.04-v1.7",
        "memory": 6111,
        "cpu": 256,
        "resourceRequirements" : [{
          "type" : "GPU",
          "value" : "1"
        }],
        "essential": true,
        "portMappings": [
          {
            "containerPort": 80,
            "protocol": "tcp"
          }
        ],
        "logConfiguration": {
            "logDriver": "awslogs",
            "options": {
                "awslogs-group": "/ecs/pytorch-s3-plugin-gpu",
                "awslogs-region": "us-east-1",
                "awslogs-stream-prefix": "imagenet",
                    "awslogs-create-group": "true"
            }
          }
        }
      ],
      "volumes": [],
      "networkMode": "bridge",
      "placementConstraints": [],
      "family": "pytorch-s3-plugin"
    }
```

3. Register the task definition. Note the revision number in the output and use it in the next step.

```
aws ecs register-task-definition --cli-input-json file://ecs-deep-learning-container-
pytorch-s3-plugin-taskdef.json
```

4. Create a task using the task definition. You need the revision identifier from the previous step.

```
aws ecs run-task --cluster ecs-pytorch-s3-plugin --task-definition pytorch-s3-plugin:1
```

5. Open the Amazon ECS console at https://console.aws.amazon.com/ecs/.

6. Select the ecs-pytorch-s3-plugin cluster.

7. On the **Cluster** page, choose **Tasks**.

8. After your task is in a RUNNING state, choose the task identifier.

9. Under **Containers**, expand the container details.

10. Under **Log Configuration**, choose **View logs in CloudWatch**. This takes you to the CloudWatch console to view the Amazon S3 plugin example logs.

For more information and additional examples, see the Amazon S3 Plugin for PyTorch repository.

Next steps

To learn inference on Amazon ECS using PyTorch with Deep Learning Containers, see PyTorch inference (p. 33).

Inference

This section shows how to run inference on AWS Deep Learning Containers for Amazon Elastic Container Service (Amazon ECS) using Apache MXNet (Incubating), PyTorch, TensorFlow, and TensorFlow 2. You can also use Elastic Inference to run inference with AWS Deep Learning Containers. For tutorials and more information on Elastic Inference, see Using AWS Deep Learning Containers with Elastic Inference on Amazon ECS.

For a complete list of Deep Learning Containers, see Deep Learning Containers Images (p. 83).

> **Note**
> MKL users: Read the AWS Deep Learning Containers Intel Math Kernel Library (MKL) Recommendations (p. 85) to get the best training or inference performance.

> **Important**
> If your account has already created the Amazon ECS service-linked role, then that role is used by default for your service unless you specify a role here. The service-linked role is required if your task definition uses the **awsvpc** network mode. The role is also required if the service is configured to use service discovery, an external deployment controller, multiple target groups, or Elastic Inference accelerators in which case you should not specify a role here. For more information, see Using Service-Linked Roles for Amazon ECS in the *Amazon ECS Developer Guide*.

Contents

- TensorFlow inference (p. 25)
- Apache MXNet (Incubating) inference (p. 29)
- PyTorch inference (p. 33)

TensorFlow inference

The following examples use a sample Docker image that adds either CPU or GPU inference scripts to Deep Learning Containers from your host machine's command line.

CPU-based inference

Use the following example to run CPU-based inference.

1. Create a file named `ecs-dlc-cpu-inference-taskdef.json` with the following contents. You can use this with either TensorFlow or TensorFlow 2. To use it with TensorFlow 2, change the Docker image to a TensorFlow 2 image and clone the r2.0 serving repository branch instead of r1.15.

```
{
  "requiresCompatibilities": [
   "EC2"
  ],
  "containerDefinitions": [{
   "command": [
    "mkdir -p /test && cd /test && git clone -b r1.15 https://github.com/
tensorflow/serving.git && tensorflow_model_server --port=8500 --rest_api_port=8501
 --model_name=saved_model_half_plus_two --model_base_path=/test/serving/
tensorflow_serving/servables/tensorflow/testdata/saved_model_half_plus_two_cpu"
   ],
```

```
  "entryPoint": [
   "sh",
   "-c"
  ],
  "name": "tensorflow-inference-container",
  "image": "763104351884.dkr.ecr.us-east-1.amazonaws.com/tensorflow-inference:1.15.0-
cpu-py36-ubuntu18.04",
  "memory": 8111,
  "cpu": 256,
  "essential": true,
  "portMappings": [{
    "hostPort": 8500,
    "protocol": "tcp",
    "containerPort": 8500
   },
   {
    "hostPort": 8501,
    "protocol": "tcp",
    "containerPort": 8501
   },
   {
    "containerPort": 80,
    "protocol": "tcp"
   }
  ],
  "logConfiguration": {
   "logDriver": "awslogs",
   "options": {
    "awslogs-group": "/ecs/tensorflow-inference-gpu",
    "awslogs-region": "us-east-1",
    "awslogs-stream-prefix": "half-plus-two",
    "awslogs-create-group": "true"
   }
  }
 }],
 "volumes": [],
 "networkMode": "bridge",
 "placementConstraints": [],
 "family": "tensorflow-inference"
}
```

2. Register the task definition. Note the revision number in the output and use it in the next step.

```
aws ecs register-task-definition --cli-input-json file://ecs-dlc-cpu-inference-
taskdef.json
```

3. Create an Amazon ECS service. When you specify the task definition, replace `revision_id` with the revision number of the task definition from the output of the previous step.

```
aws ecs create-service --cluster ecs-ec2-training-inference \
                 --service-name cli-ec2-inference-cpu \
                 --task-definition Ec2TFInference:revision_id \
                 --desired-count 1 \
                 --launch-type EC2 \
                 --scheduling-strategy="REPLICA" \
                 --region us-east-1
```

4. Verify the service and get the network endpoint by completing the following steps.

   a. Open the Amazon ECS console at https://console.aws.amazon.com/ecs/.

   b. Select the `ecs-ec2-training-inference` cluster.

   c. On the **Cluster** page, choose **Services** and then **cli-ec2-inference-cpu**.

   d. After your task is in a `RUNNING` state, choose the task identifier.

e. Under **Containers**, expand the container details.

f. Under **Name** and then **Network Bindings**, under **External Link** note the IP address for port 8501 and use it in the next step.

g. Under **Log Configuration**, choose **View logs in CloudWatch**. This takes you to the CloudWatch console to view the training progress logs.

5. To run inference, use the following command. Replace the external IP address with the external link IP address from the previous step.

```
curl -d '{"instances": [1.0, 2.0, 5.0]}' -X POST http://<External ip>:8501/v1/models/
saved_model_half_plus_two:predict
```

The following is sample output.

```
{
    "predictions": [2.5, 3.0, 4.5
    ]
}
```

> **Important**
> If you are unable to connect to the external IP address, be sure that your corporate firewall is not blocking non-standards ports, like 8501. You can try switching to a guest network to verify.

GPU-based inference

Use the following example to run GPU-based inference.

1. Create a file named `ecs-dlc-gpu-inference-taskdef.json` with the following contents. You can use this with either TensorFlow or TensorFlow 2. To use it with TensorFlow 2, change the Docker image to a TensorFlow 2 image and clone the r2.0 serving repository branch instead of r1.15.

```
{
 "requiresCompatibilities": [
  "EC2"
 ],
 "containerDefinitions": [{
  "command": [
    "mkdir -p /test && cd /test && git clone -b r1.15 https://github.com/
tensorflow/serving.git && tensorflow_model_server --port=8500 --rest_api_port=8501
 --model_name=saved_model_half_plus_two --model_base_path=/test/serving/
tensorflow_serving/servables/tensorflow/testdata/saved_model_half_plus_two_gpu"
  ],
  "entryPoint": [
   "sh",
   "-c"
  ],
  "name": "tensorflow-inference-container",
  "image": "763104351884.dkr.ecr.us-east-1.amazonaws.com/tensorflow-inference:1.15.0-
gpu-py36-cu100-ubuntu18.04",
  "memory": 8111,
  "cpu": 256,
  "resourceRequirements": [{
   "type": "GPU",
   "value": "1"
  }],
  "essential": true,
  "portMappings": [{
    "hostPort": 8500,
```

```
        "protocol": "tcp",
        "containerPort": 8500
      },
      {
        "hostPort": 8501,
        "protocol": "tcp",
        "containerPort": 8501
      },
      {
        "containerPort": 80,
        "protocol": "tcp"
      }
    ],
    "logConfiguration": {
      "logDriver": "awslogs",
      "options": {
        "awslogs-group": "/ecs/TFInference",
        "awslogs-region": "us-east-1",
        "awslogs-stream-prefix": "ecs",
        "awslogs-create-group": "true"
      }
    }
  }],
  "volumes": [],
  "networkMode": "bridge",
  "placementConstraints": [],
  "family": "TensorFlowInference"
}
```

2. Register the task definition. Note the revision number in the output and use it in the next step.

```
aws ecs register-task-definition --cli-input-json file://ecs-dlc-gpu-inference-
taskdef.json
```

3. Create an Amazon ECS service. When you specify the task definition, replace `revision_id` with the revision number of the task definition from the output of the previous step.

```
aws ecs create-service --cluster ecs-ec2-training-inference \
                       --service-name cli-ec2-inference-gpu \
                       --task-definition Ec2TFInference:revision_id \
                       --desired-count 1 \
                       --launch-type EC2 \
                       --scheduling-strategy="REPLICA" \
                       --region us-east-1
```

4. Verify the service and get the network endpoint by completing the following steps.

   a. Open the Amazon ECS console at https://console.aws.amazon.com/ecs/.

   b. Select the `ecs-ec2-training-inference` cluster.

   c. On the **Cluster** page, choose **Services** and then **cli-ec2-inference-cpu**.

   d. After your task is in a `RUNNING` state, choose the task identifier.

   e. Under **Containers**, expand the container details.

   f. Under **Name** and then **Network Bindings**, under **External Link** note the IP address for port 8501 and use it in the next step.

   g. Under **Log Configuration**, choose **View logs in CloudWatch**. This takes you to the CloudWatch console to view the training progress logs.

5. To run inference, use the following command. Replace the external IP address with the external link IP address from the previous step.

```
curl -d '{"instances": [1.0, 2.0, 5.0]}' -X POST http://<External ip>:8501/v1/models/
saved_model_half_plus_two:predict
```

The following is sample output.

```
{
    "predictions": [2.5, 3.0, 4.5
    ]
}
```

> **Important**
> If you are unable to connect to the external IP address, be sure that your corporate firewall
> is not blocking non-standards ports, like 8501. You can try switching to a guest network to
> verify.

Apache MXNet (Incubating) inference

Before you can run a task on your Amazon ECS cluster, you must register a task definition. Task
definitions are lists of containers grouped together. The following examples use a sample Docker image
that adds either CPU or GPU inference scripts to Deep Learning Containers from your host machine's
command line.

CPU-based inference

Use the following task definition to run CPU-based inference.

1. Create a file named `ecs-dlc-cpu-inference-taskdef.json` with the following contents.

```
{
 "requiresCompatibilities": [
  "EC2"
 ],
 "containerDefinitions": [{
  "command": [
   "mxnet-model-server --start --mms-config /home/model-server/config.properties
 --models squeezenet=https://s3.amazonaws.com/model-server/models/squeezenet_v1.1/
squeezenet_v1.1.model"
  ],
  "name": "mxnet-inference-container",
  "image": "763104351884.dkr.ecr.us-east-1.amazonaws.com/mxnet-inference:1.6.0-cpu-
py36-ubuntu16.04",
  "memory": 8111,
  "cpu": 256,
  "essential": true,
  "portMappings": [{
    "hostPort": 8081,
    "protocol": "tcp",
    "containerPort": 8081
   },
   {
    "hostPort": 80,
    "protocol": "tcp",
    "containerPort": 8080
   }
  ],
  "logConfiguration": {
   "logDriver": "awslogs",
   "options": {
```

```
      "awslogs-group": "/ecs/mxnet-inference-cpu",
      "awslogs-region": "us-east-1",
      "awslogs-stream-prefix": "squeezenet",
      "awslogs-create-group": "true"
    }
  }
}],
"volumes": [],
"networkMode": "bridge",
"placementConstraints": [],
"family": "mxnet-inference"
}
```

2. Register the task definition. Note the revision number in the output and use it in the next step.

```
aws ecs register-task-definition --cli-input-json file://ecs-dlc-cpu-inference-
taskdef.json
```

3. Create an Amazon ECS service. When you specify the task definition, replace `revision_id` with the revision number of the task definition from the output of the previous step.

```
aws ecs create-service --cluster ecs-ec2-training-inference \
                       --service-name cli-ec2-inference-cpu \
                       --task-definition Ec2TFInference:revision_id \
                       --desired-count 1 \
                       --launch-type EC2 \
                       --scheduling-strategy REPLICA \
                       --region us-east-1
```

4. Verify the service and get the endpoint.

    a. Open the Amazon ECS console at https://console.aws.amazon.com/ecs/.

    b. Select the `ecs-ec2-training-inference` cluster.

    c. On the **Cluster** page, choose **Services** and then **cli-ec2-inference-cpu**.

    d. After your task is in a `RUNNING` state, choose the task identifier.

    e. Under **Containers**, expand the container details.

    f. Under **Name** and then **Network Bindings**, under **External Link** note the IP address for port 8081 and use it in the next step.

    g. Under **Log Configuration**, choose **View logs in CloudWatch**. This takes you to the CloudWatch console to view the training progress logs.

5. To run inference, use the following command. Replace the `external IP` address with the external link IP address from the previous step.

```
curl -O https://s3.amazonaws.com/model-server/inputs/kitten.jpg
curl -X POST http://<External ip>/predictions/squeezenet -T kitten.jpg
```

The following is sample output.

```
[
  {
    "probability": 0.8582226634025574,
    "class": "n02124075 Egyptian cat"
  },
  {
    "probability": 0.09160050004720688,
    "class": "n02123045 tabby, tabby cat"
  },
  {
    "probability": 0.037487514317035675,
```

```
          "class": "n02123159 tiger cat"
      },
      {
          "probability": 0.0061649843119084835,
          "class": "n02128385 leopard, Panthera pardus"
      },
      {
          "probability": 0.003171598305925727,
          "class": "n02127052 lynx, catamount"
      }
  ]
```

Important

If you are unable to connect to the external IP address, be sure that your corporate firewall is not blocking non-standards ports, like 8081. You can try switching to a guest network to verify.

GPU-based inference

Use the following task definition to run GPU-based inference.

```
{
 "requiresCompatibilities": [
  "EC2"
 ],
 "containerDefinitions": [{
  "command": [
   "mxnet-model-server --start --mms-config /home/model-server/config.properties
 --models  squeezenet=https://s3.amazonaws.com/model-server/models/squeezenet_v1.1/
squeezenet_v1.1.model"
  ],
  "name": "mxnet-inference-container",
  "image": "763104351884.dkr.ecr.us-east-1.amazonaws.com/mxnet-inference:1.6.0-gpu-py36-
cu101-ubuntu16.04",
  "memory": 8111,
  "cpu": 256,
  "resourceRequirements": [{
   "type": "GPU",
   "value": "1"
  }],
  "essential": true,
  "portMappings": [{
    "hostPort": 8081,
    "protocol": "tcp",
    "containerPort": 8081
   },
   {
    "hostPort": 80,
    "protocol": "tcp",
    "containerPort": 8080
   }
  ],
  "logConfiguration": {
   "logDriver": "awslogs",
   "options": {
    "awslogs-group": "/ecs/mxnet-inference-gpu",
    "awslogs-region": "us-east-1",
    "awslogs-stream-prefix": "squeezenet",
    "awslogs-create-group": "true"
   }
  }
 }],
 "volumes": [],
```

```
 "networkMode": "bridge",
 "placementConstraints": [],
 "family": "mxnet-inference"
}
```

1. Use the following command to register the task definition. Note the output of the revision number and use it in the next step.

```
aws ecs register-task-definition --cli-input-json file://<Task definition file>
```

2. To create the service, replace the `revision_id` with the output from the previous step in the following command.

```
aws ecs create-service --cluster ecs-ec2-training-inference \
               --service-name cli-ec2-inference-gpu \
               --task-definition Ec2TFInference:<revision_id> \
               --desired-count 1 \
               --launch-type "EC2" \
               --scheduling-strategy REPLICA \
               --region us-east-1
```

3. Verify the service and get the endpoint.

   a. Open the Amazon ECS console at https://console.aws.amazon.com/ecs/.

   b. Select the `ecs-ec2-training-inference` cluster.

   c. On the **Cluster** page, choose **Services** and then **cli-ec2-inference-cpu**.

   d. After your task is in a `RUNNING` state, choose the task identifier.

   e. Under **Containers**, expand the container details.

   f. Under **Name** and then **Network Bindings**, under **External Link** note the IP address for port 8081 and use it in the next step.

   g. Under **Log Configuration**, choose **View logs in CloudWatch**. This takes you to the CloudWatch console to view the training progress logs.

4. To run inference, use the following command. Replace the `external IP` address with the external link IP address from the previous step.

```
curl -O https://s3.amazonaws.com/model-server/inputs/kitten.jpg
curl -X POST http://<External ip>/predictions/squeezenet -T kitten.jpg
```

The following is sample output.

```
[
  {
    "probability": 0.8582226634025574,
    "class": "n02124075 Egyptian cat"
  },
  {
    "probability": 0.09160050004720688,
    "class": "n02123045 tabby, tabby cat"
  },
  {
    "probability": 0.037487514317035675,
    "class": "n02123159 tiger cat"
  },
  {
    "probability": 0.0061649843119084835,
    "class": "n02128385 leopard, Panthera pardus"
  },
  {
```

```
    "probability": 0.003171598305925727,
    "class": "n02127052 lynx, catamount"
  }
]
```

Important

If you are unable to connect to the external IP address, be sure that your corporate firewall is not blocking non-standards ports, like 8081. You can try switching to a guest network to verify.

PyTorch inference

Before you can run a task on your Amazon ECS cluster, you must register a task definition. Task definitions are lists of containers grouped together. The following examples use a sample Docker image that adds either CPU or GPU inference scripts to Deep Learning Containers.

CPU-based inference

Use the following task definition to run CPU-based inference.

1. Create a file named `ecs-dlc-cpu-inference-taskdef.json` with the following contents.

```
{
    "requiresCompatibilities": [
        "EC2"
    ],
    "containerDefinitions": [{
        "command": [
            "mxnet-model-server --start --mms-config /home/model-server/
config.properties --models densenet=https://dlc-samples.s3.amazonaws.com/pytorch/multi-
model-server/densenet/densenet.mar"
        ],
        "name": "pytorch-inference-container",
        "image": "763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-inference:1.3.1-
cpu-py36-ubuntu16.04",
        "memory": 8111,
        "cpu": 256,
        "essential": true,
        "portMappings": [{
                "hostPort": 8081,
                "protocol": "tcp",
                "containerPort": 8081
            },
            {
                "hostPort": 80,
                "protocol": "tcp",
                "containerPort": 8080
            }
        ],
        "logConfiguration": {
            "logDriver": "awslogs",
            "options": {
                "awslogs-group": "/ecs/densenet-inference-cpu",
                "awslogs-region": "us-east-1",
                "awslogs-stream-prefix": "densenet",
                "awslogs-create-group": "true"
            }
        }
    }],
    "volumes": [],
    "networkMode": "bridge",
```

```
      "placementConstraints": [],
      "family": "pytorch-inference"
}
```

2. Register the task definition. Note the revision number in the output and use it in the next step.

```
aws ecs register-task-definition --cli-input-json file://ecs-dlc-cpu-inference-
taskdef.json
```

3. Create an Amazon ECS service. When you specify the task definition, replace `revision_id` with the revision number of the task definition from the output of the previous step.

```
aws ecs create-service --cluster ecs-ec2-training-inference \
                       --service-name cli-ec2-inference-cpu \
                       --task-definition Ec2PTInference:revision_id \
                       --desired-count 1 \
                       --launch-type EC2 \
                       --scheduling-strategy REPLICA \
                       --region us-east-1
```

4. Verify the service and get the network endpoint by completing the following steps.

    a. Open the Amazon ECS console at https://console.aws.amazon.com/ecs/.

    b. Select the `ecs-ec2-training-inference` cluster.

    c. On the **Cluster** page, choose **Services** and then **cli-ec2-inference-cpu**.

    d. After your task is in a `RUNNING` state, choose the task identifiter.

    e. Under **Containers**, expand the container details.

    f. Under **Name** and then **Network Bindings**, under **External Link** note the IP address for port 8081 and use it in the next step.

    g. Under **Log Configuration**, choose **View logs in CloudWatch**. This takes you to the CloudWatch console to view the training progress logs.

5. To run inference, use the following command. Replace the `external IP` address with the external link IP address from the previous step.

```
curl -O https://s3.amazonaws.com/model-server/inputs/flower.jpg
curl -X POST http://<External ip>/predictions/densenet -T flower.jpg
```

> **Important**
> If you are unable to connect to the external IP address, be sure that your corporate firewall is not blocking non-standards ports, like 8081. You can try switching to a guest network to verify.

GPU-based inference

Use the following task definition to run GPU-based inference.

```
{
    "requiresCompatibilities": [
        "EC2"
    ],
    "containerDefinitions": [{
        "command": [
            "mxnet-model-server --start --mms-config /home/model-server/config.properties
 --models densenet=https://dlc-samples.s3.amazonaws.com/pytorch/multi-model-server/
densenet/densenet.mar"
        ],
        "name": "pytorch-inference-container",
```

```
        "image": "763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-inference:1.3.1-gpu-
py36-cu101-ubuntu16.04",
        "memory": 8111,
        "cpu": 256,
        "essential": true,
        "portMappings": [{
                "hostPort": 8081,
                "protocol": "tcp",
                "containerPort": 8081
            },
            {
                "hostPort": 80,
                "protocol": "tcp",
                "containerPort": 8080
            }
        ],
        "logConfiguration": {
            "logDriver": "awslogs",
            "options": {
                "awslogs-group": "/ecs/densenet-inference-cpu",
                "awslogs-region": "us-east-1",
                "awslogs-stream-prefix": "densenet",
                "awslogs-create-group": "true"
            }
        }
    }],
    "volumes": [],
    "networkMode": "bridge",
    "placementConstraints": [],
    "family": "pytorch-inference"
}
```

1. Use the following command to register the task definition. Note the output of the revision number and use it in the next step.

```
aws ecs register-task-definition --cli-input-json file://<Task definition file>
```

2. To create the service, replace the `revision_id` with the output from the previous step in the following command.

```
aws ecs create-service --cluster ecs-ec2-training-inference \
               --service-name cli-ec2-inference-gpu \
               --task-definition Ec2PTInference:<revision_id> \
               --desired-count 1 \
               --launch-type "EC2" \
               --scheduling-strategy REPLICA \
               --region us-east-1
```

3. Verify the service and get the network endpoint by completing the following steps.

   a. Open the Amazon ECS console at https://console.aws.amazon.com/ecs/.

   b. Select the `ecs-ec2-training-inference` cluster.

   c. On the **Cluster** page, choose **Services** and then **cli-ec2-inference-cpu**.

   d. Once your task is in a `RUNNING` state, choose the task identifier.

   e. Under **Containers**, expand the container details.

   f. Under **Name** and then **Network Bindings**, under **External Link** note the IP address for port 8081 and use it in the next step.

   g. Under **Log Configuration**, choose **View logs in CloudWatch**. This takes you to the CloudWatch console to view the training progress logs.

4. To run inference, use the following command. Replace the `external IP` address with the external link IP address from the previous step.

    ```
    curl -O https://s3.amazonaws.com/model-server/inputs/flower.jpg
    curl -X POST http://<External ip>/predictions/densenet -T flower.jpg
    ```

    > **Important**
    > If you are unable to connect to the external IP address, be sure that your corporate firewall is not blocking non-standards ports, like 8081. You can try switching to a guest network to verify.

Next steps

To learn about using Custom Entrypoints with Deep Learning Containers on Amazon ECS, see .

Custom entrypoints

For some images, Deep Learning Containers use a custom entrypoint script. If you want to use your own entrypoint, you can override the entrypoint as follows.

Modify the `entryPoint` parameter in the JSON file that includes your task definition. Include the file path to your custom entry point script. An example is shown here.

```
"entryPoint":[
          "sh",
          "-c",
          "/usr/local/bin/mxnet-model-server --start --foreground --mms-config /home/
model-server/config.properties --models densenet=https://dlc-samples.s3.amazonaws.com/
pytorch/multi-model-server/densenet/densenet.mar"],
```

Amazon EKS Tutorials

This section shows how to run training and inference on AWS Deep Learning Containers for EKS using MXNet, PyTorch, and TensorFlow. Some examples cover single node or multi-node training. Inference uses only single node configurations.

Before starting the following tutorials, complete the steps in .

Contents

Amazon EKS Setup

This guide explains how to setup a deep learning environment using Amazon Elastic Kubernetes Service (Amazon EKS) and AWS Deep Learning Containers. Using Amazon EKS you can scale a production-ready environment for multiple-node training and inference with Kubernetes containers.

If you're not familiar with Kubernetes or Amazon EKS yet, that's okay. This guide and the related Amazon EKS documentation shows how to use the family of Kubernetes tools. This guide assumes that you're already familiar with your deep learning framework's multiple-node implementations and how to set up an inference server outside of containers.

A deep learning container setup on Amazon EKS consists of one or more containers, forming a cluster. You can have dedicated cluster types, such as a cluster for training and a cluster for inference. You might also want different instance types for your clusters depending on the demands of your deep learning neural networks and models.

Contents

Custom Images

Custom images are helpful if you want to load your own code or datasets and have them available on each node in your cluster. Examples are provided that use custom images. You can try them out to get started without creating your own.

- Building AWS Deep Learning Containers Custom Images (p. 84)

Licensing

To use GPU hardware, use an Amazon Machine Image that has the necessary GPU drivers. We recommend using the Amazon EKS-optimized AMI with GPU support, which is used in subsequent steps of this guide. This AMI includes non-AWS software that requires an end user license agreement (EULA). You must subscribe to the EKS-optimized AMI in the AWS Marketplace and accept the EULA before you can use the AMI in your worker node groups.

> **Important**
> To subscribe to the AMI, visit the AWS Marketplace.

Configure Security Settings

To use Amazon EKS you must have a user account that has access to several security permissions. These are set with the AWS Identity and Access Management (IAM) tool.

1. Create an IAM user or update an existing IAM user by following the steps in Creating an IAM user in your AWS account.
2. Get the credentials of this user.

   a. Open the IAM console at https://console.aws.amazon.com/iam/.

    b. Under `Users`, select your user.

    c. Select `Security Credentials`.

    d. Select `Create access key`.

    e. Download the key pair or copy the information for use later.

3. Add the following policies to your IAM user. These policies provide the required access for Amazon EKS, IAM, and Amazon Elastic Compute Cloud (Amazon EC2).

    a. Select `Permissions`.

    b. Select `Add permissions`.

    c. Select `Create policy`.

    d. From the `Create policy` window, select the `JSON` tab.

    e. Paste the following content.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "VisualEditor0",
            "Effect": "Allow",
            "Action": "eks:*",
            "Resource": "*"
        }
    ]
}
```

    f. Name the policy `EKSFullAccess` and create the policy.

    g. Navigate back to the `Grant permissions` window.

    h. Select `Attach existing policies directly`.

    i. Search for `EKSFullAccess`, and select the check box.

    j. Search for `AWSCloudFormationFullAccess`, and select the check box.

    k. Search for `AmazonEC2FullAccess`, and select the check box.

    l. Search for `IAMFullAccess`, and select the check box.

    m. Search `AmazonEC2ContainerRegistryReadOnly`, and select the check box.

    n. Search `AmazonEKS_CNI_Policy`, and select the check box.

    o. Search `AmazonS3FullAccess`, and select the check box.

    p. Accept the changes.

Gateway Node

To setup an Amazon EKS cluster, use the open source tool, `eksctl`. We recommend that you use an Amazon EC2 instance with the Deep Learning Base AMI (Ubuntu) to allocate and control your cluster. You can run these tools locally on your computer or an Amazon EC2 instance that you already have running. However, to simplify this guide we assume you're using a Deep Learning Base AMI (DLAMI) with Ubuntu 16.04. We refer to this as your gateway node.

Before you start, consider the location of your training data or where you want to run your cluster for responding to inference requests. Typically your data and cluster for training or inference should be in the same Region. Also, you spin up your gateway node in this same Region. You can follow this quick 10 minute tutorial that guides you to launch a DLAMI to use as your gateway node.

1. Login to your gateway node.

2. Install or upgrade AWS CLI. To access the required new Kubernetes features, you must have the latest version.

```
$ sudo pip install --upgrade awscli
```

3. Install eksctl by running the following commands. For more information about eksctl, see the eksctl documentation.

```
$ curl --silent \
--location "https://github.com/weaveworks/eksctl/releases/download/latest_release/
eksctl_$(uname -s)_amd64.tar.gz" \
| tar xz -C /tmp
$ sudo mv /tmp/eksctl /usr/local/bin
```

4. Install `kubectl` by following the steps in the Installing kubectl guide.

   Note
   You must use a `kubectl` version that is within one minor version difference of your Amazon EKS cluster control plane version. For example, a 1.18 `kubectl` client works with Kubernetes 1.17, 1.18 and 1.19 clusters.

5. Install `aws-iam-authenticator` by running the following commands. For more information on aws-iam-authenticator, see Installing `aws-iam-authenticator`.

```
$ curl -o aws-iam-authenticator https://amazon-eks.s3.us-
west-2.amazonaws.com/1.19.6/2021-01-05/bin/linux/amd64/aws-iam-authenticator
$ chmod +x aws-iam-authenticator
$ cp ./aws-iam-authenticator $HOME/bin/aws-iam-authenticator && export PATH=$HOME/bin:
$PATH
```

6. Run `aws configure` for the IAM user from the Security Configuration section. You are copying the IAM user's AWS Access Key, then the AWS Secret Access Key that you accessed in the IAM console and pasting these into the prompts from `aws configure`.

GPU Clusters

1. Examine the following command to create a cluster using a p3.8xlarge instance type. You must make the following modifications before you run it.

   - `name` is what you will use to manage your cluster. You can change `cluster-name` to be whatever name you like as long as there are no spaces or special characters.

   - `eks-version` is the Amazon EKS kubernetes version. For the supported Amazon EKS versions, see Available Amazon EKS Kubernetes versions.

   - `nodes` is the number of instances you want in your cluster. In this example, we're starting with three nodes.

   - `node-type` refers to instance class. You can choose a different instance class if you already know what kind will work best for your situation.

   - `timeout` and *ssh-access * can be left alone.

   - `ssh-public-key` is the name of the key that you want to use to login your worker nodes. Either use a security key you already use or create a new one but be sure to swap out the ssh-public-key with a key that was allocated for the Region you used. Note: You only need to provide the key name as seen in the 'key pairs' section of the Amazon EC2 Console.

   - `region` is the Amazon EC2 Region where the cluster will be launched. If you plan to use training data that resides in a specific Region (other than *<us-east-1>*) we recommend that you use the same Region. The ssh-public-key must have access to launch instances in this Region.

     Note
     The rest of this guide assumes *<us-east-1>* as the Region.

2. After you have made changes to the command, run it, and wait. It can take several minutes for a single node cluster, and will take even longer if you chose to create a large cluster.

```
$ eksctl create cluster <cluster-name> \
                    --version <eks-version> \
                    --nodes 3 \
                    --node-type=<p3.8xlarge> \
                    --timeout=40m \
                    --ssh-access \
                    --ssh-public-key <key_pair_name> \
                    --region <us-east-1> \
                    --zones=us-east-1a,us-east-1b,us-east-1d \
                    --auto-kubeconfig
```

You should see something similar to the following output:

```
EKS cluster "training-1" in "us-east-1" region is ready
```

3. Ideally the auto-kubeconfig should have configured your cluster. However, if you run into issues you can run the command below to set your kubeconfig. This command can also be used if you want to change your gateway node and manage your cluster from elsewhere.

```
$ aws eks --region <region> update-kubeconfig --name <cluster-name>
```

You should see something similar to the following output:

```
Added new context arn:aws:eks:us-east-1:999999999999:cluster/training-1 to /home/
ubuntu/.kube/config
```

4. If you plan to use GPU instance types, make sure to run the NVIDIA device plugin for Kubernetes on your cluster with the following command:

```
$ kubectl apply -f https://raw.githubusercontent.com/NVIDIA/k8s-device-plugin/v1.12/
nvidia-device-plugin.yml
$ kubectl create -f https://raw.githubusercontent.com/NVIDIA/k8s-device-plugin/v0.9.0/
nvidia-device-plugin.yml
```

5. Verify the GPUs available on each node in your cluster

```
$ kubectl get nodes "-o=custom-
columns=NAME:.metadata.name,GPU:.status.allocatable.nvidia\.com/gpu"
```

CPU Clusters

Refer to the previous section's discussion on using the **eksctl** command to launch a GPU cluster, and modify `node-type` to use a CPU instance type.

Habana Clusters

Refer to the previous discussion on using the **eksctl** command to launch a GPU cluster, and modify `node-type` to use an instance with Habana Gaudi accelerators, such as the DL1 instance type.

Test Your Clusters

1. You can run a **kubectl** command on the cluster to check its status. Try the command to make sure it is picking up the current cluster you want to manage.

```
$ kubectl get nodes -o wide
```

2. Take a look in ~/.kube. This directory has the kubeconfig files for the various clusters configured from your gateway node. If you browse further into the folder you can find ~/.kube/eksctl/clusters - This holds the kubeconfig file for clusters created using eksctl. This file has some details which you ideally shouldn't have to modify, since the tools are generating and updating the configurations for you, but it is good to reference when troubleshooting.

3. Verify that the cluster is active.

```
$ aws eks --region <region> describe-cluster --name <cluster-name> --query
 cluster.status
```

You should see the following output:

```
"ACTIVE"
```

4. Verify the kubectl context if you have multiple clusters set up from the same host instance. Sometimes it helps to make sure that the default context found by **kubectl** is set properly. Check this using the following command:

```
$ kubectl config get-contexts
```

5. If the context is not set as expected, fix this using the following command:

```
$ aws eks --region <region> update-kubeconfig --name <cluster-name>
```

Manage Your Clusters

When you want to control or query a cluster you can address it by the configuration file using the kubeconfig parameter. This is useful when you have more than one cluster. For example, if you have a separate cluster called "training-gpu-1" you can call the **get pods** command on it by passing the configuration file as a parameter as follows:

```
$ kubectl --kubeconfig=/home/ubuntu/.kube/eksctl/clusters/training-gpu-1 get pods
```

It is useful to note that you can run this same command without the kubeconfig parameter and it will report status on your current actively controlled cluster.

```
$ kubectl get pods
```

If you setup multiple clusters and they have yet to have the NVIDIA plugin installed, you can install it this way:

```
$ kubectl --kubeconfig=/home/ubuntu/.kube/eksctl/clusters/training-gpu-1 create -f https://
raw.githubusercontent.com/NVIDIA/k8s-device-plugin/v0.9.0/nvidia-device-plugin.yml
```

You also change the active cluster by updating the kubeconfig, passing the name of the cluster you want to manage. The following command updates kubeconfig and removes the need to use the kubeconfig parameter.

```
$ aws eks —region us-east-1 update-kubeconfig —name training-gpu-1
```

If you follow all of the examples in this guide, you will frequently switch between active clusters. This is so you can orchestrate training or inference or use different frameworks running on different clusters.

Cleanup

When you're done using the cluster, delete it to avoid incurring additional costs.

```
$ eksctl delete cluster --name=<cluster-name>
```

To delete only a pod, run the following:

```
$ kubectl delete pods <name>
```

To reset the secret for access to the cluster, run the following:

```
$ kubectl delete secret ${SECRET} -n ${NAMESPACE} || true
```

To delete a `nodegroup` attached to a cluster, run the following:

```
$ eksctl delete nodegroup --name <cluster_name>
```

To attach a `nodegroup` to a cluster, run the following:

```
$ eksctl create nodegroup
                 --cluster <cluster-name> \
                 --node-ami <ami_id> \
                 --nodes <num_nodes> \
                 --node-type=<instance_type> \
                 --timeout=40m \
                 --ssh-access \
                 --ssh-public-key <key_pair_name> \
                 --region <us-east-1> \
                 --auto-kubeconfig
```

Next steps

To learn about training and inference with Deep Learning Containers on Amazon EKS, see Amazon EKS Tutorials (p. 36).

Training

Once you've created a cluster using the steps in Amazon EKS Setup (p. 36), you can use it to run training jobs. For training, you can use either a CPU, GPU, or distributed GPU example depending on the nodes in your cluster. The topics in the following sections show how to use Apache MXNet (Incubating), PyTorch, TensorFlow, and TensorFlow 2 training examples.

Contents

CPU Training

This section is for training on CPU-based containers.

For a complete list of Deep Learning Containers, see Deep Learning Containers Images (p. 83). For tips about the best configuration settings if you're using the Intel Math Kernel Library (MKL), see AWS Deep Learning Containers Intel Math Kernel Library (MKL) Recommendations (p. 85).

Contents

Apache MXNet (Incubating) CPU training

This tutorial guides you on training with Apache MXNet (Incubating) on your single node CPU cluster.

1. Create a pod file for your cluster. A pod file will provide the instructions about what the cluster should run. This pod file will download the MXNet repository and run an MNIST example. Open **vi** or **vim** and copy and past the following content. Save this file as `mxnet.yaml`.

```
apiVersion: v1
kind: Pod
metadata:
  name: mxnet-training
spec:
  restartPolicy: OnFailure
  containers:
  - name: mxnet-training
    image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/mxnet-inference:1.6.0-cpu-py36-ubuntu16.04
    command: ["/bin/sh","-c"]
    args: ["git clone -b v1.4.x https://github.com/apache/incubator-mxnet.git &&
 python ./incubator-mxnet/example/image-classification/train_mnist.py"]
```

2. Assign the pod file to the cluster using **kubectl**.

```
$ kubectl create -f mxnet.yaml
```

3. You should see the following output:

```
pod/mxnet-training created
```

4. Check the status. The name of the job "mxnet-training" was in the mxnet.yaml file. It will now appear in the status. If you're running any other tests or have previously run something, it appears in this list. Run this several times until you see the status change to "Running".

```
$ kubectl get pods
```

You should see the following output:

```
NAME READY STATUS RESTARTS AGE
mxnet-training 0/1 Running 8 19m
```

5. Check the logs to see the training output.

```
$ kubectl logs mxnet-training
```

You should see something similar to the following output:

```
Cloning into 'incubator-mxnet'...
INFO:root:Epoch[0] Batch [0-100]    Speed: 18437.78 samples/sec    accuracy=0.777228
INFO:root:Epoch[0] Batch [100-200]   Speed: 16814.68 samples/sec    accuracy=0.907188
INFO:root:Epoch[0] Batch [200-300]   Speed: 18855.48 samples/sec    accuracy=0.926719
INFO:root:Epoch[0] Batch [300-400]   Speed: 20260.84 samples/sec    accuracy=0.938438
INFO:root:Epoch[0] Batch [400-500]   Speed: 9062.62 samples/sec    accuracy=0.938594
INFO:root:Epoch[0] Batch [500-600]   Speed: 10467.17 samples/sec    accuracy=0.945000
INFO:root:Epoch[0] Batch [600-700]   Speed: 11082.03 samples/sec    accuracy=0.954219
INFO:root:Epoch[0] Batch [700-800]   Speed: 11505.02 samples/sec    accuracy=0.956875
INFO:root:Epoch[0] Batch [800-900]   Speed: 9072.26 samples/sec    accuracy=0.955781
INFO:root:Epoch[0] Train-accuracy=0.923424
...
```

6. Check the logs to watch the training progress. You can also continue to check "**get pods**" to refresh the status. When the status changes to "`Completed`", the training job is done.

Next steps

To learn CPU-based inference on Amazon EKS using MXNet with Deep Learning Containers, see Apache MXNet (Incubating) CPU inference (p. 64).

TensorFlow CPU training

This tutorial guides you on training TensorFlow models on your single node CPU cluster.

1. Create a pod file for your cluster. A pod file will provide the instructions about what the cluster should run. This pod file will download Keras and run a Keras example. This example uses the TensorFlow framework. Open **vi** or **vim** and copy and paste the following content. Save this file as `tf.yaml`. You can use this with either TensorFlow or TensorFlow 2. To use it with TensorFlow 2, change the Docker image to a TensorFlow 2 image.

```
apiVersion: v1
kind: Pod
metadata:
  name: tensorflow-training
spec:
  restartPolicy: OnFailure
  containers:
  - name: tensorflow-training
    image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/tensorflow-inference:1.15.2-
cpu-py36-ubuntu18.04
    command: ["/bin/sh","-c"]
    args: ["git clone https://github.com/fchollet/keras.git && python /keras/examples/
mnist_cnn.py"]
```

2. Assign the pod file to the cluster using **kubectl**.

```
$ kubectl create -f tf.yaml
```

3. You should see the following output:

```
pod/tensorflow-training created
```

4. Check the status. The name of the job "tensorflow-training" was in the tf.yaml file. It will now appear in the status. If you're running any other tests or have previously run something, it appears in this list. Run this several times until you see the status change to "Running".

```
$ kubectl get pods
```

You should see the following output:

```
NAME READY STATUS RESTARTS AGE
tensorflow-training 0/1 Running 8 19m
```

5. Check the logs to see the training output.

```
$ kubectl logs tensorflow-training
```

You should see something similar to the following output:

```
Cloning into 'keras'...
Using TensorFlow backend.
Downloading data from https://s3.amazonaws.com/img-datasets/mnist.npz

    8192/11490434 [..............................] - ETA: 0s
 6479872/11490434 [===============>..............] - ETA: 0s
 8740864/11490434 [=====================>........] - ETA: 0s
11493376/11490434 [==============================] - 0s 0us/step
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
Train on 60000 samples, validate on 10000 samples
Epoch 1/12
2019-03-19 01:52:33.863598: I tensorflow/core/platform/cpu_feature_guard.cc:141] Your
 CPU supports instructions that this TensorFlow binary was not compiled to use: AVX512F
2019-03-19 01:52:33.867616: I tensorflow/core/common_runtime/process_util.cc:69]
 Creating new thread pool with default inter op setting: 2. Tune using
 inter_op_parallelism_threads for best performance.

  128/60000 [..............................] - ETA: 10:43 - loss: 2.3076 - acc: 0.0625
  256/60000 [..............................] - ETA: 5:59 - loss: 2.2528 - acc: 0.1445
  384/60000 [..............................] - ETA: 4:24 - loss: 2.2183 - acc: 0.1875
  512/60000 [..............................] - ETA: 3:35 - loss: 2.1652 - acc: 0.1953
  640/60000 [..............................] - ETA: 3:05 - loss: 2.1078 - acc: 0.2422
  ...
```

6. You can check the logs to watch the training progress. You can also continue to check "**get pods**" to refresh the status. When the status changes to "`Completed`" you will know that the training job is done.

Next steps

To learn CPU-based inference on Amazon EKS using TensorFlow with Deep Learning Containers, see .

PyTorch CPU training

This tutorial guides you on training with PyTorch on your single node CPU cluster.

1. Create a pod file for your cluster. A pod file will provide the instructions about what the cluster should run. This pod file will download the PyTorch repository and run an MNIST example. Open **vi** or **vim**, then copy and paste the following content. Save this file as `pytorch.yaml`.

```
apiVersion: v1
kind: Pod
```

```
metadata:
  name: pytorch-training
spec:
  restartPolicy: OnFailure
  containers:
  - name: pytorch-training
    image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-training:1.5.1-cpu-
py36-ubuntu16.04
    command:
      - "/bin/sh"
      - "-c"
    args:
    - "git clone https://github.com/pytorch/examples.git && python examples/mnist/
main.py --no-cuda"
    env:
    - name: OMP_NUM_THREADS
      value: "36"
    - name: KMP_AFFINITY
      value: "granularity=fine,verbose,compact,1,0"
    - name: KMP_BLOCKTIME
      value: "1"
```

2. Assign the pod file to the cluster using **kubectl**.

```
$ kubectl create -f pytorch.yaml
```

3. You should see the following output:

```
pod/pytorch-training created
```

4. Check the status. The name of the job "pytorch-training" was in the pytorch.yaml file. It will now appear in the status. If you're running any other tests or have previously run something, it appears in this list. Run this several times until you see the status change to "Running".

```
$ kubectl get pods
```

You should see the following output:

```
NAME READY STATUS RESTARTS AGE
pytorch-training 0/1 Running 8 19m
```

5. Check the logs to see the training output.

```
$ kubectl logs pytorch-training
```

You should see something similar to the following output:

```
Cloning into 'examples'...
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to ../data/
MNIST/raw/train-images-idx3-ubyte.gz
9920512it [00:00, 40133996.38it/s]
Extracting ../data/MNIST/raw/train-images-idx3-ubyte.gz to ../data/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to ../data/
MNIST/raw/train-labels-idx1-ubyte.gz
Extracting ../data/MNIST/raw/train-labels-idx1-ubyte.gz to ../data/MNIST/raw
32768it [00:00, 831315.84it/s]
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to ../data/
MNIST/raw/t10k-images-idx3-ubyte.gz
1654784it [00:00, 13019129.43it/s]
Extracting ../data/MNIST/raw/t10k-images-idx3-ubyte.gz to ../data/MNIST/raw
```

```
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to ../data/
MNIST/raw/t10k-labels-idx1-ubyte.gz
8192it [00:00, 337197.38it/s]
Extracting ../data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ../data/MNIST/raw
Processing...
Done!
Train Epoch: 1 [0/60000 (0%)]    Loss: 2.300039
Train Epoch: 1 [640/60000 (1%)]    Loss: 2.213470
Train Epoch: 1 [1280/60000 (2%)]    Loss: 2.170460
Train Epoch: 1 [1920/60000 (3%)]    Loss: 2.076699
Train Epoch: 1 [2560/60000 (4%)]    Loss: 1.868078
Train Epoch: 1 [3200/60000 (5%)]    Loss: 1.414199
Train Epoch: 1 [3840/60000 (6%)]    Loss: 1.000870
```

6. Check the logs to watch the training progress. You can also continue to check "**get pods**" to refresh the status. When the status changes to "`Completed`" you will know that the training job is done.

See EKS Cleanup for information on cleaning up a cluster after you're done using it.

Amazon S3 Plugin for PyTorch

Deep Learning Containers include a plugin that enables you to use data from an Amazon S3 bucket for PyTorch training.

1. To begin using the Amazon S3 plugin on Amazon EKS, check to make sure that your cluster instances have full access to Amazon S3. Create an IAM role that grants Amazon S3 access to an Amazon EC2 instance and attach the role to your instance. You can use the AmazonS3FullAccess or AmazonS3ReadOnlyAccess policies.

2. Set up your `AWS_REGION` environment variable with the region of your choice.

```
export AWS_REGION=us-east-1
```

3. Create a pod file for your cluster. A pod file will provide the instructions about what the cluster should run. This pod file will use the PyTorch Amazon S3 plugin to access an example Amazon S3 dataset.

> **Note**
> Your CPU cluster should use `c5.12xlarge` nodes or greater for this example.

Open **vi** or **vim**, then copy and paste the following content. Save this file as `s3plugin.yaml`.

```
apiVersion: v1
kind: Pod
metadata:
  name: pytorch-s3-plugin
spec:
  restartPolicy: OnFailure
  containers:
  - name: pytorch-s3-plugin
    image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-training:1.8.1-cpu-
py36-ubuntu18.04-v1.6
    command:
      - "/bin/sh"
      - "-c"
    args:
    - "git clone https://github.com/aws/amazon-s3-plugin-for-pytorch.git && python
 amazon-s3-plugin-for-pytorch/examples/s3_imagenet_example.py"
    env:
    - name: OMP_NUM_THREADS
      value: "36"
    - name: KMP_AFFINITY
```

```
        value: "granularity=fine,verbose,compact,1,0"
    - name: KMP_BLOCKTIME
      value: "1"
```

4. Assign the pod file to the cluster using **kubectl**.

    ```
    $ kubectl create -f s3plugin.yaml
    ```

5. Check the status. The name of the job `pytorch-s3-plugin` that was specified in the
    `s3plugin.yaml` file will now appear alongside the status information. You can run the following
    command several times until you see the status change to "`Running`."

    ```
    $ kubectl get pods
    ```

    You should see the following output:

    ```
    NAME READY STATUS RESTARTS AGE
    pytorch-s3-plugin 0/1 Running 8 19m
    ```

6. Check the logs to see more details.

    ```
    $ kubectl logs pytorch-s3-plugin
    ```

For more information, see the Amazon S3 Plugin for PyTorch repository.

Next steps

To learn CPU-based inference on Amazon EKS using PyTorch with Deep Learning Containers, see PyTorch
CPU inference (p. 69).

GPU Training

This section is for training on GPU-based clusters.

For a complete list of Deep Learning Containers, refer to Deep Learning Containers Images (p. 83). For
tips about the best configuration settings if you're using the Intel Math Kernel Library (MKL), see AWS
Deep Learning Containers Intel Math Kernel Library (MKL) Recommendations (p. 85).

Contents

Apache MXNet (Incubating) GPU training

This tutorial guides you on training with Apache MXNet (Incubating) on your single node GPU cluster.

1. Create a pod file for your cluster. A pod file will provide the instructions about what the cluster
    should run. This pod file will download the MXNet repository and run an MNIST example. Open **vi** or
    vim and copy and past the following content. Save this file as `mxnet.yaml`.

    ```
    apiVersion: v1
    ```

```
kind: Pod
metadata:
  name: mxnet-training
spec:
  restartPolicy: OnFailure
  containers:
  - name: mxnet-training
    image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/mxnet-training:1.6.0-gpu-py36-
cu101-ubuntu16.04
    command: ["/bin/sh","-c"]
    args: ["git clone -b v1.4.x https://github.com/apache/incubator-mxnet.git &&
 python ./incubator-mxnet/example/image-classification/train_mnist.py"]
```

2. Assign the pod file to the cluster using **kubectl**.

```
$ kubectl create -f mxnet.yaml
```

3. You should see the following output:

```
pod/mxnet-training created
```

4. Check the status. The name of the job "tensorflow-training" was in the tf.yaml file. It will now appear in the status. If you're running any other tests or have previously run something, it will appear in this list. Run this several times until you see the status change to "`Running`".

```
$ kubectl get pods
```

You should see the following output:

```
NAME READY STATUS RESTARTS AGE
mxnet-training 0/1 Running 8 19m
```

5. Check the logs to see the training output.

```
$ kubectl logs mxnet-training
```

You should see something similar to the following output:

```
Cloning into 'incubator-mxnet'...
INFO:root:Epoch[0] Batch [0-100]    Speed: 18437.78 samples/sec    accuracy=0.777228
INFO:root:Epoch[0] Batch [100-200]   Speed: 16814.68 samples/sec    accuracy=0.907188
INFO:root:Epoch[0] Batch [200-300]   Speed: 18855.48 samples/sec    accuracy=0.926719
INFO:root:Epoch[0] Batch [300-400]   Speed: 20260.84 samples/sec    accuracy=0.938438
INFO:root:Epoch[0] Batch [400-500]   Speed: 9062.62 samples/sec    accuracy=0.938594
INFO:root:Epoch[0] Batch [500-600]   Speed: 10467.17 samples/sec    accuracy=0.945000
INFO:root:Epoch[0] Batch [600-700]   Speed: 11082.03 samples/sec    accuracy=0.954219
INFO:root:Epoch[0] Batch [700-800]   Speed: 11505.02 samples/sec    accuracy=0.956875
INFO:root:Epoch[0] Batch [800-900]   Speed: 9072.26 samples/sec    accuracy=0.955781
INFO:root:Epoch[0] Train-accuracy=0.923424
...
```

6. Check the logs to watch the training progress. You can also continue to check "**get pods**" to refresh the status. When the status changes to "`Completed`", the training job is done.

Next steps

To learn GPU-based inference on Amazon EKS using MXNet with Deep Learning Containers, see Apache MXNet (Incubating) GPU inference (p. 71).

TensorFlow GPU training

This tutorial guides you on training TensorFlow models on your single node GPU cluster.

1. Create a pod file for your cluster. A pod file will provide the instructions about what the cluster should run. This pod file will download Keras and run a Keras example. This example uses the TensorFlow framework. Open **vi** or **vim** and copy and past the following content. Save this file as `tf.yaml`. You can use this with either TensorFlow or TensorFlow 2. To use it with TensorFlow 2, change the Docker image to a TensorFlow 2 image.

```
apiVersion: v1
kind: Pod
metadata:
  name: tensorflow-training
spec:
  restartPolicy: OnFailure
  containers:
  - name: tensorflow-training
    image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/tensorflow-training:1.15.2-gpu-
py37-cu100-ubuntu18.04
    command: ["/bin/sh","-c"]
    args: ["git clone https://github.com/fchollet/keras.git && python /keras/examples/
mnist_cnn.py"]
    resources:
      limits:
        nvidia.com/gpu: 1
```

2. Assign the pod file to the cluster using **kubectl**.

```
$ kubectl create -f tf.yaml
```

3. You should see the following output:

```
pod/tensorflow-training created
```

4. Check the status. The name of the job "tensorflow-training" was in the tf.yaml file. It will now appear in the status. If you're running any other tests or have previously run something, it appears in this list. Run this several times until you see the status change to "`Running`".

```
$ kubectl get pods
```

You should see the following output:

```
NAME READY STATUS RESTARTS AGE
tensorflow-training 0/1 Running 8 19m
```

5. Check the logs to see the training output.

```
$ kubectl logs tensorflow-training
```

You should see something similar to the following output:

```
Cloning into 'keras'...
Using TensorFlow backend.
Downloading data from https://s3.amazonaws.com/img-datasets/mnist.npz

    8192/11490434 [..............................] - ETA: 0s
 6479872/11490434 [===============>..............] - ETA: 0s
```

```
 8740864/11490434 [====================>........] - ETA: 0s
11493376/11490434 [==============================] - 0s 0us/step
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
Train on 60000 samples, validate on 10000 samples
Epoch 1/12
2019-03-19 01:52:33.863598: I tensorflow/core/platform/cpu_feature_guard.cc:141] Your
 CPU supports instructions that this TensorFlow binary was not compiled to use: AVX512F
2019-03-19 01:52:33.867616: I tensorflow/core/common_runtime/process_util.cc:69]
 Creating new thread pool with default inter op setting: 2. Tune using
 inter_op_parallelism_threads for best performance.

  128/60000 [..............................] - ETA: 10:43 - loss: 2.3076 - acc: 0.0625
  256/60000 [..............................] - ETA: 5:59 - loss: 2.2528 - acc: 0.1445
  384/60000 [..............................] - ETA: 4:24 - loss: 2.2183 - acc: 0.1875
  512/60000 [..............................] - ETA: 3:35 - loss: 2.1652 - acc: 0.1953
  640/60000 [..............................] - ETA: 3:05 - loss: 2.1078 - acc: 0.2422
  ...
```

6. Check the logs to watch the training progress. You can also continue to check "**get pods**" to refresh the status. When the status changes to "`Completed`", the training job is done.

Next steps

To learn GPU-based inference on Amazon EKS using TensorFlow with Deep Learning Containers, see TensorFlow GPU inference (p. 73).

PyTorch GPU training

This tutorial guides you on training with PyTorch on your single node GPU cluster.

1. Create a pod file for your cluster. A pod file will provide the instructions about what the cluster should run. This pod file will download the PyTorch repository and run an MNIST example. Open **vi** or **vim**, then copy and paste the following content. Save this file as `pytorch.yaml`.

```
apiVersion: v1
kind: Pod
metadata:
  name: pytorch-training
spec:
  restartPolicy: OnFailure
  containers:
  - name: pytorch-training
    image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-training:1.5.1-gpu-
py36-cu101-ubuntu16.04
    command:
      - "/bin/sh"
      - "-c"
    args:
    - "git clone https://github.com/pytorch/examples.git && python examples/mnist/
main.py --no-cuda"
    env:
    - name: OMP_NUM_THREADS
      value: "36"
    - name: KMP_AFFINITY
      value: "granularity=fine,verbose,compact,1,0"
    - name: KMP_BLOCKTIME
      value: "1"
```

2. Assign the pod file to the cluster using **kubectl**.

```
$ kubectl create -f pytorch.yaml
```

3. You should see the following output:

```
pod/pytorch-training created
```

4. Check the status. The name of the job "pytorch-training" was in the pytorch.yaml file. It will now appear in the status. If you're running any other tests or have previously run something, it appears in this list. Run this several times until you see the status change to "Running".

```
$ kubectl get pods
```

You should see the following output:

```
NAME READY STATUS RESTARTS AGE
pytorch-training 0/1 Running 8 19m
```

5. Check the logs to see the training output.

```
$ kubectl logs pytorch-training
```

You should see something similar to the following output:

```
Cloning into 'examples'...
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to ../data/
MNIST/raw/train-images-idx3-ubyte.gz
9920512it [00:00, 40133996.38it/s]
Extracting ../data/MNIST/raw/train-images-idx3-ubyte.gz to ../data/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to ../data/
MNIST/raw/train-labels-idx1-ubyte.gz
Extracting ../data/MNIST/raw/train-labels-idx1-ubyte.gz to ../data/MNIST/raw
32768it [00:00, 831315.84it/s]
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to ../data/
MNIST/raw/t10k-images-idx3-ubyte.gz
1654784it [00:00, 13019129.43it/s]
Extracting ../data/MNIST/raw/t10k-images-idx3-ubyte.gz to ../data/MNIST/raw
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to ../data/
MNIST/raw/t10k-labels-idx1-ubyte.gz
8192it [00:00, 337197.38it/s]
Extracting ../data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ../data/MNIST/raw
Processing...
Done!
Train Epoch: 1 [0/60000 (0%)]    Loss: 2.300039
Train Epoch: 1 [640/60000 (1%)]    Loss: 2.213470
Train Epoch: 1 [1280/60000 (2%)]    Loss: 2.170460
Train Epoch: 1 [1920/60000 (3%)]    Loss: 2.076699
Train Epoch: 1 [2560/60000 (4%)]    Loss: 1.868078
Train Epoch: 1 [3200/60000 (5%)]    Loss: 1.414199
Train Epoch: 1 [3840/60000 (6%)]    Loss: 1.000870
```

6. Check the logs to watch the training progress. You can also continue to check "**get pods**" to refresh the status. When the status changes to "Completed", the training job is done.

See EKS Cleanup for information on cleaning up a cluster after you're done using it.

Next steps

To learn GPU-based inference on Amazon EKS using PyTorch with Deep Learning Containers, see PyTorch GPU inference (p. 77).

Amazon S3 Plugin for PyTorch

Deep Learning Containers include a plugin that enables you to use data from an Amazon S3 bucket for PyTorch training.

1. To begin using the Amazon S3 plugin on Amazon EKS, check to make sure that your cluster instances have full access to Amazon S3. Create an IAM role that grants Amazon S3 access to an Amazon EC2 instance and attach the role to your instance. You can use the AmazonS3FullAccess or AmazonS3ReadOnlyAccess policies.

2. Set up your `AWS_REGION` environment variable with the region of your choice.

    ```
    export AWS_REGION=us-east-1
    ```

3. Create a pod file for your cluster. A pod file will provide the instructions about what the cluster should run. This pod file will use the PyTorch Amazon S3 plugin to access an example Amazon S3 dataset.

    Note
    Your GPU cluster should use `p3.8xlarge` nodes or greater for this example.

    Open **vi** or **vim**, then copy and paste the following content. Save this file as `s3plugin.yaml`.

    ```
    apiVersion: v1
    kind: Pod
    metadata:
      name: pytorch-s3-plugin
    spec:
      restartPolicy: OnFailure
      containers:
      - name: pytorch-s3-plugin
        image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-training:1.8.1-gpu-
    py36-cu111-ubuntu18.04-v1.7
        command:
          - "/bin/sh"
          - "-c"
        args:
        - "git clone https://github.com/aws/amazon-s3-plugin-for-pytorch.git && python
     amazon-s3-plugin-for-pytorch/examples/s3_imagenet_example.py"
        env:
        - name: OMP_NUM_THREADS
          value: "36"
        - name: KMP_AFFINITY
          value: "granularity=fine,verbose,compact,1,0"
        - name: KMP_BLOCKTIME
          value: "1"
    ```

4. Assign the pod file to the cluster using **kubectl**.

    ```
    $ kubectl create -f s3plugin.yaml
    ```

5. Check the status. The name of the job `pytorch-s3-plugin` that was specified in the `s3plugin.yaml` file will now appear alongside the status information. You can run the following command several times until you see the status change to "Running."

    ```
    $ kubectl get pods
    ```

You should see the following output:

```
NAME READY STATUS RESTARTS AGE
pytorch-s3-plugin 0/1 Running 8 19m
```

6. Check the logs to see more details.

```
$ kubectl logs pytorch-s3-plugin
```

For more information, see the Amazon S3 Plugin for PyTorch repository.

Distributed GPU Training

This section is for running distributed training on multi-node GPU clusters.

For a complete list of Deep Learning Containers, refer to Deep Learning Containers Images (p. 83).

Contents

Set up your cluster for distributed training

To run distributed training on EKS, you need the following components installed on your cluster.

- The default installation of Kubeflow with required components, such as PyTorch operators, TensorFlow operators, and the NVIDIA plugin.
- Apache MXNet and MPI operators.

Download and run the script to install the required components in the cluster.

```
$ wget -O install_kubeflow.sh https://raw.githubusercontent.com/aws/deep-learning-
containers/master/test/dlc_tests/eks/eks_manifest_templates/kubeflow/install_kubeflow.sh
$ chmod +x install_kubeflow.sh
$ ./install_kubeflow.sh <EKS_CLUSTER_NAME> <AWS_REGION>
```

Apache MXNet (Incubating) distributed GPU training

This tutorial shows how to run distributed training with Apache MXNet (Incubating) on your multi-node GPU cluster using Parameter Server. To run MXNet distributed training on EKS, you use the Kubernetes MXNet-operator named MXJob. It provides a custom resource that makes it easy to run distributed or non-distributed MXNet jobs (training and tuning) on Kubernetes. This operator is installed in the previous setup step.

Using a Custom Resource Definition (CRD) gives users the ability to create and manage MX Jobs just like builtin K8s resources. Verify that the MXNet custom resource is installed.

```
$ kubectl get crd
```

The output should include `mxjobs.kubeflow.org`.

Running MNIST distributed training with parameter server example

Create a pod file(mx_job_dist.yaml) for your job according to the available cluster configuration and job to run. There are 3 jobModes you need to specify: Scheduler, Server and Worker. You can specify how many pods you want to spawn with the field replicas. The instance type of the Scheduler, Server, and Worker will be of the type specified at cluster creation.

- Scheduler: There is only one scheduler. The role of the scheduler is to set up the cluster. This includes waiting for messages that each node has come up and which port the node is listening on. The scheduler then lets all processes know about every other node in the cluster, so that they can communicate with each other.
- Server: There can be multiple servers which store the model's parameters, and communicate with workers. A server may or may not be co-located with the worker processes.
- Worker: A worker node actually performs training on a batch of training samples. Before processing each batch, the workers pull weights from servers. The workers also send gradients to the servers after each batch. Depending on the workload for training a model, it might not be a good idea to run multiple worker processes on the same machine.
- Provide container image you want to use with the field image.
- You can provide `restartPolicy` from one of the Always, OnFailure and Never. It determines whether pods will be restarted when they exit or not.
- Provide container image you want to use with the field image.

1. To create a MXJob template, modify the following code block according to your requirements and save it in a file named `mx_job_dist.yaml`.

```
apiVersion: "kubeflow.org/v1beta1"
kind: "MXJob"
metadata:
  name: <JOB_NAME>
spec:
  jobMode: MXTrain
  mxReplicaSpecs:
    Scheduler:
      replicas: 1
      restartPolicy: Never
      template:
        spec:
          containers:
            - name: mxnet
              image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/aws-samples-mxnet-training:1.8.0-gpu-py37-cu110-ubuntu16.04-example
    Server:
      replicas: <NUM_SERVERS>
      restartPolicy: Never
      template:
        spec:
          containers:
            - name: mxnet
              image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/aws-samples-mxnet-training:1.8.0-gpu-py37-cu110-ubuntu16.04-example
    Worker:
      replicas: <NUM_WORKERS>
      restartPolicy: Never
      template:
        spec:
          containers:
            - name: mxnet
```

```
                  image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/aws-samples-mxnet-
training:1.8.0-gpu-py37-cu110-ubuntu16.04-example
              command:
              - "python"
              args:
              - "/incubator-mxnet/example/image-classification/train_mnist.py"
              - "--num-epochs"
              - <EPOCHS>
              - "--num-layers"
              - <LAYERS>
              - "--kv-store"
              - "dist_device_sync"
              - "--gpus"
              - <GPUS>
              resources:
                limits:
                  nvidia.com/gpu: <GPU_LIMIT>
```

2. Run distributed training job with the pod file you just created.

```
$ # Create a job by defining MXJob
kubectl create -f mx_job_dist.yaml
```

3. List the running jobs.

```
$ kubectl get mxjobs
```

4. To get status of a running job, run the following. Replace the JOB variable with whatever the job's name is.

```
$ JOB=<JOB_NAME>
kubectl get mxjobs $JOB -o yaml
```

The output should be similar to the following:

```
apiVersion: kubeflow.org/v1beta1
kind: MXJob
metadata:
  creationTimestamp: "2020-07-23T16:38:41Z"
  generation: 8
  name: kubeflow-mxnet-gpu-dist-job-3910
  namespace: mxnet-multi-node-training-3910
  resourceVersion: "688398"
  selfLink: /apis/kubeflow.org/v1beta1/namespaces/mxnet-multi-node-training-3910/
mxjobs/kubeflow-mxnet-gpu-dist-job-3910
spec:
  cleanPodPolicy: All
  jobMode: MXTrain
  mxReplicaSpecs:
    Scheduler:
      replicas: 1
      restartPolicy: Never
      template:
        metadata:
          creationTimestamp: null
        spec:
          containers:
          - image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/aws-samples-mxnet-
training:1.8.0-gpu-py37-cu110-ubuntu16.04-example
            name: mxnet
            ports:
            - containerPort: 9091
```

```
            name: mxjob-port
          resources: {}
    Server:
      replicas: 2
      restartPolicy: Never
      template:
        metadata:
          creationTimestamp: null
        spec:
          containers:
          - image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/aws-samples-mxnet-
training:1.8.0-gpu-py37-cu110-ubuntu16.04-example
            name: mxnet
            ports:
            - containerPort: 9091
              name: mxjob-port
            resources: {}
    Worker:
      replicas: 3
      restartPolicy: Never
      template:
        metadata:
          creationTimestamp: null
        spec:
          containers:
          - args:
            - /incubator-mxnet/example/image-classification/train_mnist.py
            - --num-epochs
            - "20"
            - --num-layers
            - "2"
            - --kv-store
            - dist_device_sync
            - --gpus
            - "0"
            command:
            - python
            image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/aws-samples-mxnet-
training:1.8.0-gpu-py37-cu110-ubuntu16.04-example
            name: mxnet
            ports:
            - containerPort: 9091
              name: mxjob-port
            resources:
              limits:
                nvidia.com/gpu: "1"
status:
  conditions:
  - lastTransitionTime: "2020-07-23T16:38:41Z"
    lastUpdateTime: "2020-07-23T16:38:41Z"
    message: MXJob kubeflow-mxnet-gpu-dist-job-3910 is created.
    reason: MXJobCreated
    status: "True"
    type: Created
  - lastTransitionTime: "2020-07-23T16:38:41Z"
    lastUpdateTime: "2020-07-23T16:40:50Z"
    message: MXJob kubeflow-mxnet-gpu-dist-job-3910 is running.
    reason: MXJobRunning
    status: "True"
    type: Running
  mxReplicaStatuses:
    Scheduler:
      active: 1
    Server:
      active: 2
    Worker:
```

```
        active: 3
   startTime: "2020-07-23T16:40:50Z"
```

> **Note**
> Status provides information about the state of the resources.
> Phase - Indicates the phase of a job and will be one of Creating, Running, CleanUp, Failed,
> or Done.
> State - Provides the overall status of the job and will be one of Running, Succeeded, or
> Failed.

5. If you want to delete a job, change directories to where you launched the job and run the following:

```
$ kubectl delete -f mx_job_dist.yaml
```

Apache MXNet (Incubating) with Horovod distributed GPU training

This tutorial shows how to setup distributed training of Apache MXNet (Incubating) models on your multi-node GPU cluster that uses Horovod. It uses an example image that already has a training script included, and it uses a 3-node cluster with `node-type=p3.8xlarge`. This tutorial runs the Horovod example script for MXNet on an MNIST model.

1. Verify that the MPIJob custom resource is installed.

```
$ kubectl get crd
```

The output should include `mpijobs.kubeflow.org`.

2. Create a MPI Job template and define the number of nodes (replicas) and number of GPUs each node has (gpusPerReplica). Modify the following code block according to your requirements and save it in a file named `mx-mnist-horovod-job.yaml`.

```
apiVersion: kubeflow.org/v1alpha2
kind: MPIJob
metadata:
  name: <JOB_NAME>
spec:
  slotsPerWorker: 1
  cleanPodPolicy: Running
  mpiReplicaSpecs:
    Launcher:
      replicas: 1
      template:
        spec:
          containers:
          - image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/aws-samples-mxnet-
training:1.8.0-gpu-py37-cu110-ubuntu16.04-example
            name: <JOB_NAME>
            args:
            - --epochs
            - "10"
            - --lr
            - "0.001"
            command:
            - mpirun
            - -mca
            - btl_tcp_if_exclude
            - lo
            - -mca
            - pml
            - ob1
```

```
                    - -mca
                    - btl
                    - ^openib
                    - --bind-to
                    - none
                    - -map-by
                    - slot
                    - -x
                    - LD_LIBRARY_PATH
                    - -x
                    - PATH
                    - -x
                    - NCCL_SOCKET_IFNAME=eth0
                    - -x
                    - NCCL_DEBUG=INFO
                    - -x
                    - MXNET_CUDNN_AUTOTUNE_DEFAULT=0
                    - python
                    - /horovod/examples/mxnet_mnist.py
        Worker:
          replicas: <NUM_WORKERS>
          template:
            spec:
              containers:
                - image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/aws-samples-mxnet-
        training:1.8.0-gpu-py37-cu110-ubuntu16.04-example
                  name: mpi-worker
                  resources:
                    limits:
                      nvidia.com/gpu: <GPUS>
```

3. Run the distributed training job with the pod file you just created.

    ```
    $ kubectl create -f mx-mnist-horovod-job.yaml
    ```

4. Check the status. The name of the job appears in the status. If you're running any other tests or
    have previously run something, it appears in this list. Run this several times until you see the status
    change to "Running".

    ```
    $ kubectl get pods -o wide
    ```

    You should see something similar to the following output:

    ```
    NAME                                          READY    STATUS    RESTARTS AGE

    mxnet-mnist-horovod-job-716-launcher-4wc7f    1/1      Running   0        31s
    mxnet-mnist-horovod-job-716-worker-0          1/1      Running   0        31s
    mxnet-mnist-horovod-job-716-worker-1          1/1      Running   0        31s
    mxnet-mnist-horovod-job-716-worker-2          1/1      Running   0        31s
    ```

5. Based on the name of the launcher pod above, check the logs to see the training output.

    ```
    $ kubectl logs -f --tail 10 <LAUNCHER_POD_NAME>
    ```

6. You can check the logs to watch the training progress. You can also continue to check "get pods" to
    refresh the status. When the status changes to "Completed" you will know that the training job is
    done.

7. To clean up and rerun a job:

    ```
    $ kubectl delete -f mx-mnist-horovod-job.yaml
    ```

To learn GPU-based inference on Amazon EKS using MXNet with Deep Learning Containers, see Apache MXNet (Incubating) GPU inference (p. 71).

TensorFlow with Horovod distributed GPU training

This tutorial shows how to setup distributed training of TensorFlow models on your multi-node GPU cluster that uses Horovod. It uses an example image that already has a training script included, and it uses a 3-node cluster with `node-type=p3.16xlarge`. You can use this tutorial with either TensorFlow or TensorFlow 2. To use it with TensorFlow 2, change the Docker image to a TensorFlow 2 image.

1. Verify that the MPIJob custom resource is installed.

   ```
   $ kubectl get crd
   ```

   The output should include `mpijobs.kubeflow.org`.

2. Create a MPI Job template and define the number of nodes (replicas) and number of GPUs each node has (gpusPerReplica). Modify the following code block according to your requirements and save it in a file named `tf-resnet50-horovod-job.yaml`.

   ```
   apiVersion: kubeflow.org/v1alpha2
   kind: MPIJob
   metadata:
     name: <JOB_NAME>
   spec:
     slotsPerWorker: 1
     cleanPodPolicy: Running
     mpiReplicaSpecs:
       Launcher:
         replicas: 1
         template:
           spec:
             containers:
             - image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/aws-samples-
   tensorflow-training:1.15.5-gpu-py37-cu100-ubuntu18.04-example
               name: <JOB_NAME>
               command:
                - mpirun
                - -mca
                - btl_tcp_if_exclude
                - lo
                - -mca
                - pml
                - ob1
                - -mca
                - btl
                - ^openib
                - --bind-to
                - none
                - -map-by
                - slot
                - -x
                - LD_LIBRARY_PATH
                - -x
                - PATH
                - -x
                - NCCL_SOCKET_IFNAME=eth0
                - -x
                - NCCL_DEBUG=INFO
                - -x
                - MXNET_CUDNN_AUTOTUNE_DEFAULT=0
   ```

```
                - python
                - /deep-learning-models/models/resnet/tensorflow/
train_imagenet_resnet_hvd.py
              args:
                - --num_epochs
                - "10"
                - --synthetic
    Worker:
      replicas: <NUM_WORKERS>
      template:
        spec:
          containers:
          - image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/aws-samples-tensorflow-
training:1.15.5-gpu-py37-cu100-ubuntu18.04-example
            name: tensorflow-worker
            resources:
              limits:
                nvidia.com/gpu: <GPUS>
```

3. Run the distributed training job with the pod file you just created.

    ```
    $ kubectl create -f tf-resnet50-horovod-job.yaml
    ```

4. Check the status. The name of the job appears in the status. If you're running any other tests or have previously run other tests, they appear in this list. Run this several times until you see the status change to "Running".

    ```
    $ kubectl get pods -o wide
    ```

    You should see something similar to the following output:

    ```
    NAME                                         READY    STATUS    RESTARTS AGE

    tf-resnet50-horovod-job-1794-launcher-9wbsg  1/1      Running   0        31s
    tf-resnet50-horovod-job-1794-worker-0        1/1      Running   0        31s
    tf-resnet50-horovod-job-1794-worker-1        1/1      Running   0        31s
    tf-resnet50-horovod-job-1794-worker-2        1/1      Running   0        31s
    ```

5. Based on the name of the launcher pod above, check the logs to see the training output.

    ```
    $ kubectl logs -f --tail 10 <LAUNCHER_POD_NAME>
    ```

6. You can check the logs to watch the training progress. You can also continue to check "get pods" to refresh the status. When the status changes to "Completed" you will know that the training job is done.

7. To clean up and rerun a job:

    ```
    $ kubectl delete -f tf-resnet50-horovod-job.yaml
    ```

Next steps

To learn GPU-based inference on Amazon EKS using TensorFlow with Deep Learning Containers, see TensorFlow GPU inference (p. 73).

PyTorch distributed GPU training

This tutorial will guide you on distributed training with PyTorch on your multi-node GPU cluster. It uses Gloo as the backend.

1. Verify that the PyTorch custom resource is installed.

```
$ kubectl get crd
```

The output should include `pytorchjobs.kubeflow.org`.

2. Ensure that the NVIDIA plugin `daemonset` is running.

```
$ kubectl get daemonset -n kubeflow
```

The output should should look similar to the following.

```
NAME                             DESIRED CURRENT READY UP-TO-DATE AVAILABLE NODE SELECTOR
 AGE
nvidia-device-plugin-daemonset   3        3       3     3          3         <none>  35h
```

3. Use the following text to create a gloo-based distributed data parallel job. Save it in a file named `distributed.yaml`.

```
apiVersion: kubeflow.org/v1
kind: PyTorchJob
metadata:
  name: "kubeflow-pytorch-gpu-dist-job"
spec:
  pytorchReplicaSpecs:
    Master:
      replicas: 1
      restartPolicy: OnFailure
      template:
        spec:
          containers:
            - name: "pytorch"
              image: "763104351884.dkr.ecr.us-east-1.amazonaws.com/aws-samples-pytorch-training:1.7.1-gpu-py36-cu110-ubuntu18.04-example"
              args:
                - "--backend"
                - "gloo"
                - "--epochs"
                - "5"
    Worker:
      replicas: 2
      restartPolicy: OnFailure
      template:
        spec:
          containers:
            - name: "pytorch"
              image: "763104351884.dkr.ecr.us-east-1.amazonaws.com/aws-samples-pytorch-training:1.7.1-gpu-py36-cu110-ubuntu18.04-example"
              args:
                - "--backend"
                - "gloo"
                - "--epochs"
                - "5"
              resources:
                limits:
                  nvidia.com/gpu: 1
```

4. Run a distributed training job with the pod file you just created.

```
$ kubectl create -f distributed.yaml
```

5. You can check the status of the job using the following:

```
$ kubectl logs kubeflow-pytorch-gpu-dist-job
```

To view logs continuously, use:

```
$ kubectl logs -f <pod>
```

See EKS Cleanup for information on cleaning up a cluster after you're done using it.

Amazon S3 Plugin for PyTorch

Deep Learning Containers include a plugin that enables you to use data from an Amazon S3 bucket for PyTorch training. See the Amazon EKS Amazon S3 Plugin for PyTorch GPU guide to get started.

For more information and additional examples, see the Amazon S3 Plugin for PyTorch repository.

See EKS Cleanup for information on cleaning up a cluster after you're done using it.

Next steps

To learn GPU-based inference on Amazon EKS using PyTorch with Deep Learning Containers, see PyTorch GPU inference (p. 77).

Inference

Once you've created a cluster using the steps in Amazon EKS Setup (p. 36), you can use it to run inference jobs. For inference, you can use either a CPU or GPU example depending on the nodes in your cluster. Inference supports only single node configurations. The following topics show how to run inference with AWS Deep Learning Containers on EKS using Apache MXNet (Incubating), PyTorch, TensorFlow, and TensorFlow 2.

Contents

CPU Inference

This section guides you on running inference on Deep Learning Containers for EKS CPU clusters using Apache MXNet (Incubating), PyTorch, TensorFlow, and TensorFlow 2.

For a complete list of Deep Learning Containers, see Available Deep Learning Containers Images.

> **Note**
> If you're using MKL, see AWS Deep Learning Containers Intel Math Kernel Library (MKL) Recommendations (p. 85) to get the best training or inference performance.

Contents

Apache MXNet (Incubating) CPU inference

In this tutorial, you create a Kubernetes Service and a Deployment to run CPU inference with MXNet. The Kubernetes Service exposes a process and its ports. When you create a Kubernetes Service, you can specify the kind of Service you want using `ServiceTypes`. The default `ServiceType` is `ClusterIP`. The Deployment is responsible for ensuring that a certain number of pods is always up and running.

1. Create the namespace. You may need to change the kubeconfig to point to the right cluster. Verify that you have setup a "training-cpu-1" or change this to your CPU cluster's config. For more information on setting up your cluster, see Amazon EKS Setup (p. 36).

```
$ NAMESPACE=mx-inference; kubectl —kubeconfig=/home/ubuntu/.kube/eksctl/clusters/
training-cpu-1 create namespace ${NAMESPACE}
```

2. (Optional step when using public models.) Set up your model at a network location that is mountable, like in Amazon S3. For information on how to upload a trained model to S3, see TensorFlow CPU inference (p. 65). Apply the secret to your namespace. For more information on secrets, see the Kubernetes Secrets documentation.

```
$ kubectl —n ${NAMESPACE} apply -f secret.yaml
```

3. Create a file named `mx_inference.yaml` with the following content. This example file specifies the model, MXNet inference image used, and the location of the model. This example uses a public model, so you don't need to modify it.

```
---
kind: Service
apiVersion: v1
metadata:
  name: squeezenet-service
  labels:
    app: squeezenet-service
spec:
  ports:
  - port: 8080
    targetPort: mms
  selector:
    app: squeezenet-service
---
kind: Deployment
apiVersion: apps/v1
metadata:
  name: squeezenet-service
  labels:
    app: squeezenet-service
spec:
  replicas: 1
  selector:
    matchLabels:
      app: squeezenet-service
  template:
    metadata:
      labels:
        app: squeezenet-service
    spec:
      containers:
      - name: squeezenet-service
        image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/mxnet-inference:1.6.0-cpu-
py36-ubuntu16.04
        args:
        - mxnet-model-server
        - --start
```

```
        - --mms-config /home/model-server/config.properties
        - --models squeezenet=https://s3.amazonaws.com/model-server/model_archive_1.0/
squeezenet_v1.1.mar
        ports:
        - name: mms
          containerPort: 8080
        - name: mms-management
          containerPort: 8081
        imagePullPolicy: IfNotPresent
```

4. Apply the configuration to a new pod in the previously defined namespace.

```
$ kubectl -n ${NAMESPACE} apply -f mx_inference.yaml
```

Your output should be similar to the following:

```
service/squeezenet-service created
deployment.apps/squeezenet-service created
```

5. Check the status of the pod.

```
$ kubectl get pods -n ${NAMESPACE}
```

Repeat the status check until you see the following "RUNNING" state:

```
NAME                        READY     STATUS      RESTARTS    AGE
squeezenet-service-xvw1     1/1       Running     0           3m
```

6. To further describe the pod, run the following:

```
$ kubectl describe pod <pod_name> -n ${NAMESPACE}
```

7. Because the serviceType here is ClusterIP, you can forward the port from your container to your host machine using the following command:

```
$ kubectl port-forward -n ${NAMESPACE} `kubectl get pods -n ${NAMESPACE} --
selector=app=squeezenet-service -o jsonpath='{.items[0].metadata.name}'` 8080:8080 &
```

8. Download an image of a kitten.

```
$ curl -O https://s3.amazonaws.com/model-server/inputs/kitten.jpg
```

9. Run inference on the model using the image of the kitten:

```
$ curl -X POST http://127.0.0.1:8080/predictions/squeezenet -T kitten.jpg
```

TensorFlow CPU inference

In this tutorial, you create a Kubernetes Service and a Deployment to run CPU inference with TensorFlow. The Kubernetes Service exposes a process and its ports. When you create a Kubernetes Service, you can specify the kind of Service you want using `ServiceTypes`. The default `ServiceType` is `ClusterIP`. The Deployment is responsible for ensuring that a certain number of pods is always up and running.

1. Create the namespace. You may need to change the kubeconfig to point to the right cluster. Verify that you have setup a "training-cpu-1" or change this to your CPU cluster's config. For more information on setting up your cluster, see Amazon EKS Setup (p. 36).

```
$ NAMESPACE=tf-inference; kubectl —kubeconfig=/home/ubuntu/.kube/eksctl/clusters/
training-cpu-1 create namespace ${NAMESPACE}
```

2. Models served for inference can be retrieved in different ways, such as using shared volumes and Amazon S3. Because the Kubernetes Service requires access to Amazon S3 and Amazon ECR, you must store your AWS credentials as a Kubernetes secret. For the purpose of this example, use S3 to store and fetch trained models.

    Verify your AWS credentials. They must have S3 write access.

```
$ cat ~/.aws/credentials
```

3. The output will be similar to the following:

```
$ [default]
aws_access_key_id = YOURACCESSKEYID
aws_secret_access_key = YOURSECRETACCESSKEY
```

4. Encode the credentials using base64.

    Encode the access key first.

```
$ echo -n 'YOURACCESSKEYID' | base64
```

    Encode the secret access key next.

```
$ echo -n 'YOURSECRETACCESSKEY' | base64
```

    Your output should look similar to the following:

```
$ echo -n 'YOURACCESSKEYID' | base64
RkFLRUFXU0FDQ0VTU0tFWUlE
$ echo -n 'YOURSECRETACCESSKEY' | base64
RkFLRUFXU1NFQ1JFVEFDQ0VTU0tFWQ==
```

5. Create a file named `secret.yaml` with the following content in your home directory. This file is used to store the secret.

```
apiVersion: v1
kind: Secret
metadata:
  name: aws-s3-secret
type: Opaque
data:
  AWS_ACCESS_KEY_ID: YOURACCESSKEYID
  AWS_SECRET_ACCESS_KEY: YOURSECRETACCESSKEY
```

6. Apply the secret to your namespace.

```
$ kubectl -n ${NAMESPACE} apply -f secret.yaml
```

7. Clone the tensorflow-serving repository.

```
$ git clone https://github.com/tensorflow/serving/
$ cd serving/tensorflow_serving/servables/tensorflow/testdata/
```

8. Sync the pretrained `saved_model_half_plus_two_cpu` model to your S3 bucket.

```
$ aws s3 sync saved_model_half_plus_two_cpu s3://<your_s3_bucket>/
saved_model_half_plus_two
```

9. Create a file named `tf_inference.yaml` with the following content. Update `--model_base_path` to use your S3 bucket. You can use this with either TensorFlow or TensorFlow 2. To use it with TensorFlow 2, change the Docker image to a TensorFlow 2 image.

```
---
  kind: Service
  apiVersion: v1
  metadata:
    name: half-plus-two
    labels:
      app: half-plus-two
  spec:
    ports:
    - name: http-tf-serving
      port: 8500
      targetPort: 8500
    - name: grpc-tf-serving
      port: 9000
      targetPort: 9000
    selector:
      app: half-plus-two
      role: master
    type: ClusterIP
---
  kind: Deployment
  apiVersion: apps/v1
  metadata:
    name: half-plus-two
    labels:
      app: half-plus-two
      role: master
  spec:
    replicas: 1
    selector:
      matchLabels:
        app: half-plus-two
        role: master
    template:
      metadata:
        labels:
          app: half-plus-two
          role: master
      spec:
        containers:
        - name: half-plus-two
          image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/tensorflow-
inference:1.15.0-cpu-py36-ubuntu18.04
          command:
          - /usr/bin/tensorflow_model_server
          args:
          - --port=9000
          - --rest_api_port=8500
          - --model_name=saved_model_half_plus_two
          - --model_base_path=s3://tensorflow-trained-models/saved_model_half_plus_two
          ports:
          - containerPort: 8500
          - containerPort: 9000
          imagePullPolicy: IfNotPresent
          env:
          - name: AWS_ACCESS_KEY_ID
```

```
              valueFrom:
                secretKeyRef:
                  key: AWS_ACCESS_KEY_ID
                  name: aws-s3-secret
          - name: AWS_SECRET_ACCESS_KEY
            valueFrom:
              secretKeyRef:
                key: AWS_SECRET_ACCESS_KEY
                name: aws-s3-secret
          - name: AWS_REGION
            value: us-east-1
          - name: S3_USE_HTTPS
            value: "true"
          - name: S3_VERIFY_SSL
            value: "true"
          - name: S3_ENDPOINT
            value: s3.us-east-1.amazonaws.com
```

10. Apply the configuration to a new pod in the previously defined namespace.

```
$ kubectl -n ${NAMESPACE} apply -f tf_inference.yaml
```

Your output should be similar to the following:

```
service/half-plus-two created
deployment.apps/half-plus-two created
```

11. Check the status of the pod.

```
$ kubectl get pods -n ${NAMESPACE}
```

Repeat the status check until you see the following "RUNNING" state:

```
NAME                   READY    STATUS    RESTARTS   AGE
half-plus-two-vmwp9   1/1        Running    0          3m
```

12. To further describe the pod, you can run:

```
$ kubectl describe pod <pod_name> -n ${NAMESPACE}
```

13. Because the serviceType is ClusterIP, you can forward the port from your container to your host machine.

```
$ kubectl port-forward -n ${NAMESPACE} `kubectl get pods -n ${NAMESPACE} --
selector=app=half-plus-two -o jsonpath='{.items[0].metadata.name}'` 8500:8500 &
```

14. Place the following json string in a file named `half_plus_two_input.json`

```
{"instances": [1.0, 2.0, 5.0]}
```

15. Run inference on the model.

```
$ curl -d @half_plus_two_input.json -X POST http://localhost:8500/v1/models/
saved_model_half_plus_two_cpu:predict
```

Your output should look like the following:

```
{
```

```
     "predictions": [2.5, 3.0, 4.5
     ]
}
```

PyTorch CPU inference

In this approach, you create a Kubernetes Service and a Deployment to run CPU inference with PyTorch. The Kubernetes Service exposes a process and its ports. When you create a Kubernetes Service, you can specify the kind of Service you want using `ServiceTypes`. The default `ServiceType` is `ClusterIP`. The Deployment is responsible for ensuring that a certain number of pods is always up and running.

1. Create the namespace. You may need to change the kubeconfig to point to the right cluster. Verify that you have setup a "training-cpu-1" or change this to your CPU cluster's config. For more information on setting up your cluster, see Amazon EKS Setup (p. 36).

    ```
    $ NAMESPACE=pt-inference; kubectl create namespace ${NAMESPACE}
    ```

2. (Optional step when using public models.) Setup your model at a network location that is mountable, like in Amazon S3. For information on how to upload a trained model to S3, see TensorFlow CPU inference (p. 65). Apply the secret to your namespace. For more information on secrets, see the Kubernetes Secrets documentation.

    ```
    $ kubectl -n ${NAMESPACE} apply -f secret.yaml
    ```

3. Create a file named `pt_inference.yaml` with the following content. This example file specifies the model, PyTorch inference image used, and the location of the model. This example uses a public model, so you don't need to modify it.

    ```
    ---
    kind: Service
    apiVersion: v1
    metadata:
      name: densenet-service
      labels:
        app: densenet-service
    spec:
      ports:
      - port: 8080
        targetPort: mms
      selector:
        app: densenet-service
    ---
    kind: Deployment
    apiVersion: apps/v1
    metadata:
      name: densenet-service
      labels:
        app: densenet-service
    spec:
      replicas: 1
      selector:
        matchLabels:
          app: densenet-service
      template:
        metadata:
          labels:
            app: densenet-service
        spec:
          containers:
          - name: densenet-service
    ```

```
        image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-inference:1.3.1-
cpu-py36-ubuntu16.04
        args:
        - mxnet-model-server
        - --start
        - --mms-config /home/model-server/config.properties
        - --models densenet=https://dlc-samples.s3.amazonaws.com/pytorch/multi-model-
server/densenet/densenet.mar
        ports:
        - name: mms
          containerPort: 8080
        - name: mms-management
          containerPort: 8081
        imagePullPolicy: IfNotPresent
```

4. Apply the configuration to a new pod in the previously defined namespace.

```
$ kubectl -n ${NAMESPACE} apply -f pt_inference.yaml
```

Your output should be similar to the following:

```
service/densenet-service created
deployment.apps/densenet-service created
```

5. Check the status of the pod and wait for the pod to be in "RUNNING" state:

```
$ kubectl get pods -n ${NAMESPACE} -w
```

Your output should be similar to the following:

```
NAME                      READY      STATUS     RESTARTS    AGE
densenet-service-xvw1     1/1        Running    0           3m
```

6. To further describe the pod, run the following:

```
$ kubectl describe pod <pod_name> -n ${NAMESPACE}
```

7. Because the serviceType here is ClusterIP, you can forward the port from your container to your host machine.

```
$ kubectl port-forward -n ${NAMESPACE} `kubectl get pods -n ${NAMESPACE} --
selector=app=densenet-service -o jsonpath='{.items[0].metadata.name}'` 8080:8080 &
```

8. With your server started, you can now run inference from a different window using the following:

```
$ curl -O https://s3.amazonaws.com/model-server/inputs/flower.jpg
curl -X POST http://127.0.0.1:8080/predictions/densenet -T flower.jpg
```

See EKS Cleanup for information on cleaning up a cluster after you're done using it.

Next steps

To learn about using Custom Entrypoints with Deep Learning Containers on Amazon EKS, see Custom Entrypoints (p. 79).

GPU Inference

This section shows how to run inference on Deep Learning Containers for EKS GPU clusters using Apache MXNet (Incubating), PyTorch, TensorFlow, and TensorFlow 2.

For a complete list of Deep Learning Containers, see Available Deep Learning Containers Images.

> **Note**
> MKL users: read the AWS Deep Learning Containers Intel Math Kernel Library (MKL) Recommendations (p. 85) to get the best training or inference performance.

Contents

Apache MXNet (Incubating) GPU inference

In this approach, you create a Kubernetes Service and a Deployment. The Kubernetes Service exposes a process and its ports. When you create a Kubernetes Service, you can specify the kind of Service you want using `ServiceTypes`. The default `ServiceType` is `ClusterIP`. The Deployment is responsible for ensuring that a certain number of pods is always up and running.

1. For GPU-base inference, install the NVIDIA device plugin for Kubernetes:

```
$ kubectl apply -f https://raw.githubusercontent.com/NVIDIA/k8s-device-plugin/v1.12/
nvidia-device-plugin.yml
```

2. Verify that the nvidia-device-plugin-daemonset is running correctly.

```
$ kubectl get daemonset -n kube-system
```

The output will be similar to the following:

```
NAME                                DESIRED    CURRENT    READY    UP-TO-DATE    AVAILABLE
 NODE SELECTOR    AGE
aws-node                            3          3          3        3             3
 <none>          6d
kube-proxy                          3          3          3        3             3
 <none>          6d
nvidia-device-plugin-daemonset      3          3          3        3             3
 <none>          57s
```

3. Create the namespace. You might need to change the kubeconfig to point to the right cluster. Verify that you have setup a "training-gpu-1" or change this to your GPU cluster's config. For more information on setting up your cluster, see Amazon EKS Setup (p. 36).

```
$ NAMESPACE=mx-inference; kubectl —kubeconfig=/home/ubuntu/.kube/eksctl/clusters/
training-gpu-1 create namespace ${NAMESPACE}
```

4. (Optional step when using public models.) Setup your model at a network location that is mountable e.g., in S3. Refer to the steps to upload a trained model to S3 mentioned in the section Inference with TensorFlow. Apply the secret to your namespace. For more information on secrets, see the Kubernetes Secrets documentation.

```
$ kubectl -n ${NAMESPACE} apply -f secret.yaml
```

5. Create the file `mx_inference.yaml`. Use the contents of the next code block as its content.

```yaml
---
kind: Service
apiVersion: v1
metadata:
  name: squeezenet-service
  labels:
    app: squeezenet-service
spec:
  ports:
  - port: 8080
    targetPort: mms
  selector:
    app: squeezenet-service
---
kind: Deployment
apiVersion: apps/v1
metadata:
  name: squeezenet-service
  labels:
    app: squeezenet-service
spec:
  replicas: 1
  selector:
    matchLabels:
      app: squeezenet-service
  template:
    metadata:
      labels:
        app: squeezenet-service
    spec:
      containers:
      - name: squeezenet-service
        image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/mxnet-inference:1.6.0-gpu-py36-cu101-ubuntu16.04
        args:
        - mxnet-model-server
        - --start
        - --mms-config /home/model-server/config.properties
        - --models squeezenet=https://s3.amazonaws.com/model-server/model_archive_1.0/squeezenet_v1.1.mar
        ports:
        - name: mms
          containerPort: 8080
        - name: mms-management
          containerPort: 8081
        imagePullPolicy: IfNotPresent
        resources:
          limits:
            cpu: 4
            memory: 4Gi
            nvidia.com/gpu: 1
          requests:
            cpu: "1"
            memory: 1Gi
```

6. Apply the configuration to a new pod in the previously defined namespace:

```
$ kubectl -n ${NAMESPACE} apply -f mx_inference.yaml
```

Your output should be similar to the following:

```
service/squeezenet-service created
deployment.apps/squeezenet-service created
```

7. Check status of the pod and wait for the pod to be in "RUNNING" state:

```
$ kubectl get pods -n ${NAMESPACE}
```

8. Repeat the check status step until you see the following "RUNNING" state:

```
NAME                      READY     STATUS     RESTARTS     AGE
squeezenet-service-xvw1   1/1       Running    0            3m
```

9. To further describe the pod, you can run:

```
$ kubectl describe pod <pod_name> -n ${NAMESPACE}
```

10. Since the serviceType here is ClusterIP, you can forward the port from your container to your host machine (the ampersand runs this in the background):

```
$ kubectl port-forward -n ${NAMESPACE} `kubectl get pods -n ${NAMESPACE} --
selector=app=squeezenet-service -o jsonpath='{.items[0].metadata.name}'` 8080:8080 &
```

11. Download an image of a kitten:

```
$ curl -O https://s3.amazonaws.com/model-server/inputs/kitten.jpg
```

12. Run inference on the model:

```
$ curl -X POST http://127.0.0.1:8080/predictions/squeezenet -T kitten.jpg
```

TensorFlow GPU inference

In this approach, you create a Kubernetes Service and a Deployment. The Kubernetes Service exposes a process and its ports. When you create a Kubernetes Service, you can specify the kind of Service you want using ServiceTypes. The default ServiceType is ClusterIP. The Deployment is responsible for ensuring that a certain number of pods is always up and running.

1. For GPU-base inference, install the NVIDIA device plugin for Kubernetes:

```
$ kubectl apply -f https://raw.githubusercontent.com/NVIDIA/k8s-device-plugin/v1.12/
nvidia-device-plugin.yml
```

2. Verify that the nvidia-device-plugin-daemonset is running correctly.

```
$ kubectl get daemonset -n kube-system
```

The output will be similar to the following:

```
NAME                     DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE
 NODE SELECTOR   AGE
aws-node                 3         3         3       3            3
 <none>         6d
kube-proxy               3         3         3       3            3
 <none>         6d
```

```
nvidia-device-plugin-daemonset   3         3         3         3         3
  <none>           57s
```

3. Create the namespace. You might need to change the kubeconfig to point to the right cluster. Verify that you have setup a "training-gpu-1" or change this to your GPU cluster's config. For more information on setting up your cluster, see .

```
$ NAMESPACE=tf-inference; kubectl —kubeconfig=/home/ubuntu/.kube/eksctl/clusters/
training-gpu-1 create namespace ${NAMESPACE}
```

4. Models served for inference can be retrieved in different ways e.g., using shared volumes, S3 etc. Since the service will require access to S3 and ECR, you must store your AWS credentials as a Kubernetes secret. For the purpose of this example, you will use S3 to store and fetch trained models.

   Check your AWS credentials. These must have S3 write access.

   ```
   $ cat ~/.aws/credentials
   ```

5. The output will be something similar to the following:

   ```
   $ [default]
   aws_access_key_id = FAKEAWSACCESSKEYID
   aws_secret_access_key = FAKEAWSSECRETACCESSKEY
   ```

6. Encode the credentials using base64. Encode the access key first.

   ```
   $ echo —n 'FAKEAWSACCESSKEYID' | base64
   ```

   Encode the secret access key next.

   ```
   $ echo —n 'FAKEAWSSECRETACCESSKEYID' | base64
   ```

   Your output should look similar to the following:

   ```
   $ echo —n 'FAKEAWSACCESSKEYID' | base64
   RkFLRUFXU0FDQ0VTU0tFWUlE
   $ echo —n 'FAKEAWSSECRETACCESSKEY' | base64
   RkFLRUFXU1NFQ1JFVEFDQ0VTU0tFWQ==
   ```

7. Create a yaml file to store the secret. Save it as secret.yaml in your home directory.

   ```
   apiVersion: v1
   kind: Secret
   metadata:
     name: aws-s3-secret
   type: Opaque
   data:
     AWS_ACCESS_KEY_ID: RkFLRUFXU0FDQ0VTU0tFWUlE
     AWS_SECRET_ACCESS_KEY: RkFLRUFXU1NFQ1JFVEFDQ0VTU0tFWQ==
   ```

8. Apply the secret to your namespace:

   ```
   $ kubectl —n ${NAMESPACE} apply -f secret.yaml
   ```

9. In this example, you will clone the tensorflow-serving repository and sync a pretrained model to an S3 bucket. The following sample names the bucket `tensorflow-serving-models`. It also syncs a saved model to an S3 bucket called `saved_model_half_plus_two_gpu`.

```
$ git clone https://github.com/tensorflow/serving/
$ cd serving/tensorflow_serving/servables/tensorflow/testdata/
```

10. Sync the CPU model.

```
$ aws s3 sync saved_model_half_plus_two_gpu s3://<your_s3_bucket>/
saved_model_half_plus_two_gpu
```

11. Create the file `tf_inference.yaml`. Use the contents of the next code block as its content, and update `--model_base_path` to use your S3 bucket. You can use this with either TensorFlow or TensorFlow 2. To use it with TensorFlow 2, change the Docker image to a TensorFlow 2 image.

```
---
  kind: Service
  apiVersion: v1
  metadata:
    name: half-plus-two
    labels:
      app: half-plus-two
  spec:
    ports:
    - name: http-tf-serving
      port: 8500
      targetPort: 8500
    - name: grpc-tf-serving
      port: 9000
      targetPort: 9000
    selector:
      app: half-plus-two
      role: master
    type: ClusterIP
---
  kind: Deployment
  apiVersion: apps/v1
  metadata:
    name: half-plus-two
    labels:
      app: half-plus-two
      role: master
  spec:
    replicas: 1
    selector:
      matchLabels:
        app: half-plus-two
        role: master
    template:
      metadata:
        labels:
          app: half-plus-two
          role: master
      spec:
        containers:
        - name: half-plus-two
          image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/tensorflow-
inference:1.15.0-gpu-py36-cu100-ubuntu18.04
          command:
          - /usr/bin/tensorflow_model_server
          args:
          - --port=9000
          - --rest_api_port=8500
          - --model_name=saved_model_half_plus_two_gpu
```

```
            - --model_base_path=s3://tensorflow-trained-models/
saved_model_half_plus_two_gpu
          ports:
          - containerPort: 8500
          - containerPort: 9000
          imagePullPolicy: IfNotPresent
          env:
          - name: AWS_ACCESS_KEY_ID
            valueFrom:
              secretKeyRef:
                key: AWS_ACCESS_KEY_ID
                name: aws-s3-secret
          - name: AWS_SECRET_ACCESS_KEY
            valueFrom:
              secretKeyRef:
                key: AWS_SECRET_ACCESS_KEY
                name: aws-s3-secret
          - name: AWS_REGION
            value: us-east-1
          - name: S3_USE_HTTPS
            value: "true"
          - name: S3_VERIFY_SSL
            value: "true"
          - name: S3_ENDPOINT
            value: s3.us-east-1.amazonaws.com
            resources:
              limits:
                cpu: 4
                memory: 4Gi
                nvidia.com/gpu: 1
              requests:
                cpu: "1"
                memory: 1Gi
```

12. Apply the configuration to a new pod in the previously defined namespace:

```
$ kubectl -n ${NAMESPACE} apply -f tf_inference.yaml
```

Your output should be similar to the following:

```
service/half-plus-two created
deployment.apps/half-plus-two created
```

13. Check status of the pod and wait for the pod to be in "RUNNING" state:

```
$ kubectl get pods -n ${NAMESPACE}
```

14. Repeat the check status step until you see the following "RUNNING" state:

```
NAME                    READY    STATUS    RESTARTS    AGE
half-plus-two-vmwp9   1/1       Running    0            3m
```

15. To further describe the pod, you can run:

```
$ kubectl describe pod <pod_name> -n ${NAMESPACE}
```

16. Since the serviceType here is ClusterIP, you can forward the port from your container to your host machine (the ampersand runs this in the background):

```
$ kubectl port-forward -n ${NAMESPACE} `kubectl get pods -n ${NAMESPACE} --
selector=app=half-plus-two -o jsonpath='{.items[0].metadata.name}'` 8500:8500 &
```

17. Place the following json string in a file called `half_plus_two_input.json`

```
{"instances": [1.0, 2.0, 5.0]}
```

18. Run inference on the model:

```
$ curl -d @half_plus_two_input.json -X POST http://localhost:8500/v1/models/
saved_model_half_plus_two_cpu:predict
```

The expected output is as follows:

```
{
    "predictions": [2.5, 3.0, 4.5
    ]
}
```

PyTorch GPU inference

In this approach, you create a Kubernetes Service and a Deployment. The Kubernetes Service exposes a process and its ports. When you create a Kubernetes Service, you can specify the kind of Service you want using `ServiceTypes`. The default `ServiceType` is `ClusterIP`. The Deployment is responsible for ensuring that a certain number of pods is always up and running.

1. For GPU-base inference, install the NVIDIA device plugin for Kubernetes.

```
$ kubectl apply -f https://raw.githubusercontent.com/NVIDIA/k8s-device-plugin/v1.12/
nvidia-device-plugin.yml
```

2. Verify that the nvidia-device-plugin-daemonset is running correctly.

```
$ kubectl get daemonset -n kube-system
```

The output will be similar to the following.

```
NAME                           DESIRED    CURRENT    READY    UP-TO-DATE    AVAILABLE
 NODE SELECTOR     AGE
aws-node                       3          3          3        3             3
 <none>          6d
kube-proxy                     3          3          3        3             3
 <none>          6d
nvidia-device-plugin-daemonset 3          3          3        3             3
 <none>          57s
```

3. Create the namespace.

```
$ NAMESPACE=pt-inference; kubectl create namespace ${NAMESPACE}
```

4. (Optional step when using public models.) Setup your model at a network location that is mountable e.g., in S3. Refer to the steps to upload a trained model to S3 mentioned in the section Inference with TensorFlow. Apply the secret to your namespace. For more information on secrets, see the Kubernetes Secrets documentation.

```
$ kubectl -n ${NAMESPACE} apply -f secret.yaml
```

5. Create the file `pt_inference.yaml`. Use the contents of the next code block as its content.

```
---
kind: Service
apiVersion: v1
metadata:
  name: densenet-service
  labels:
    app: densenet-service
spec:
  ports:
  - port: 8080
    targetPort: mms
  selector:
    app: densenet-service
---
kind: Deployment
apiVersion: apps/v1
metadata:
  name: densenet-service
  labels:
    app: densenet-service
spec:
  replicas: 1
  selector:
    matchLabels:
      app: densenet-service
  template:
    metadata:
      labels:
        app: densenet-service
    spec:
      containers:
      - name: densenet-service
        image: "763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-inference:1.3.1-
gpu-py36-cu101-ubuntu16.04"
        args:
        - mxnet-model-server
        - --start
        - --mms-config /home/model-server/config.properties
        - --models densenet=https://dlc-samples.s3.amazonaws.com/pytorch/multi-model-
server/densenet/densenet.mar
        ports:
        - name: mms
          containerPort: 8080
        - name: mms-management
          containerPort: 8081
        imagePullPolicy: IfNotPresent
        resources:
          limits:
            cpu: 4
            memory: 4Gi
            nvidia.com/gpu: 1
          requests:
            cpu: "1"
            memory: 1Gi
```

6. Apply the configuration to a new pod in the previously defined namespace.

```
$ kubectl -n ${NAMESPACE} apply -f pt_inference.yaml
```

Your output should be similar to the following:

```
service/densenet-service created
deployment.apps/densenet-service created
```

7. Check status of the pod and wait for the pod to be in "RUNNING" state.

```
$ kubectl get pods -n ${NAMESPACE}
```

Your output should be similar to the following:

```
NAME                    READY      STATUS     RESTARTS    AGE
densenet-service-xvw1   1/1        Running    0           3m
```

8. To further describe the pod, you can run:

```
$ kubectl describe pod <pod_name> -n ${NAMESPACE}
```

9. Since the serviceType here is ClusterIP, you can forward the port from your container to your host machine (the ampersand runs this in the background).

```
$ kubectl port-forward -n ${NAMESPACE} `kubectl get pods -n ${NAMESPACE} --
selector=app=densenet-service -o jsonpath='{.items[0].metadata.name}'` 8080:8080 &
```

10. With your server started, you can now run inference from a different window.

```
$ curl -O https://s3.amazonaws.com/model-server/inputs/flower.jpg
curl -X POST http://127.0.0.1:8080/predictions/densenet -T flower.jpg
```

See EKS Cleanup for information on cleaning up a cluster after you're done using it.

Next steps

To learn about using Custom Entrypoints with Deep Learning Containers on Amazon EKS, see Custom Entrypoints (p. 79).

Custom Entrypoints

For some images, AWS containers use a custom entrypoint script. If you want to use your own entrypoint, you can override the entrypoint as follows.

Update the `command` parameter in your pod file. Replace the `args` parameters with your custom entrypoint script.

```
---
apiVersion: v1
kind: Pod
metadata:
  name: pytorch-multi-model-server-densenet
spec:
  restartPolicy: OnFailure
  containers:
  - name: pytorch-multi-model-server-densenet
    image: 763104351884.dkr.ecr.us-east-1.amazonaws.com/pytorch-inference:1.2.0-cpu-py36-
ubuntu16.04
    command:
      - "/bin/sh"
```

```
      - "-c"
    args:
      - /usr/local/bin/mxnet-model-server
      - --start
      - --mms-config /home/model-server/config.properties
      - --models densenet="https://dlc-samples.s3.amazonaws.com/pytorch/multi-model-server/
densenet/densenet.mar"
```

command is the Kubernetes field name for `entrypoint`. Refer to this table of Kubernetes field names for more information.

If the EKS cluster has expired IAM permissions to access the ECR repository holding the image, or you are using kubectl from a different user than the one that created the cluster, you will receive the following error.

```
error: unable to recognize "job.yaml": Unauthorized
```

To address this issue, you need to refresh the IAM tokens. Run the following script.

```
set -ex

AWS_ACCOUNT=${AWS_ACCOUNT}
AWS_REGION=us-east-1
DOCKER_REGISTRY_SERVER=https://${AWS_ACCOUNT}.dkr.ecr.${AWS_REGION}.amazonaws.com
DOCKER_USER=AWS
DOCKER_PASSWORD=`aws ecr get-login --region ${AWS_REGION} --registry-ids ${AWS_ACCOUNT} |
 cut -d' ' -f6`
kubectl delete secret aws-registry || true
kubectl create secret docker-registry aws-registry \
--docker-server=$DOCKER_REGISTRY_SERVER \
--docker-username=$DOCKER_USER \
--docker-password=$DOCKER_PASSWORD
kubectl patch serviceaccount default -p '{"imagePullSecrets":[{"name":"aws-registry"}]}'
```

Append the following under `spec` in your pod file.

```
imagePullSecrets:
    - name: aws-registry
```

Troubleshooting AWS Deep Learning Containers on EKS

The following are common errors that might be returned in the command line when using AWS Deep Learning Containers on an Amazon EKS cluster. Each error is followed by a solution to the error.

Troubleshooting

Topics

Setup Errors

The following errors might be returned when setting up Deep Learning Containers on your Amazon EKS cluster.

- **Error: registry `kubeflow` does not exist**

```
$ ks pkg install kubeflow/tf-serving
  ERROR registry 'kubeflow' does not exist
```

To solve this error, run the following command.

```
ks registry add kubeflow github.com/google/kubefl ow/tree/master/kubeflow
```

- **Error: context deadline exceeded**

```
$ eksctl create cluster <args>
  [#] waiting for CloudFormation stack "eksctl-training-cluster-1-nodegroup-ng-8c4c94bc"
 to reach "CREATE_COMPLETE" status: RequestCanceled: waiter context canceled
  caused by: context deadline exceeded
```

To solve this error, verify that you have not exceeded capacity for your account. You can also try to create your cluster in a different region.

- **Error: The connection to the server localhost:8080 was refused**

```
$ kubectl get nodes
  The connection to the server localhost:8080 was refused - did you specify the right
 host or port?
```

To solve this error, copy the cluster to the Kubernetes configuration by running the following.

```
cp ~/.kube/eksctl/clusters/<cluster-name> ~/.kube/config
```

- **Error: handle object: patching object from cluster: merging object with existing state: Unauthorized**

```
$ ks apply default
  ERROR handle object: patching object from cluster: merging object with existing state:
 Unauthorized
```

This error is due to a concurrency issue that can occur when multiple users with different authorization or credentials credentials try to start jobs on the same cluster. Verify that you are starting a job on the correct cluster.

- **Error: Could not create app; directory '/home/ubuntu/kubeflow-tf-hvd' already exists**

```
$ APP_NAME=kubeflow-tf-hvd; ks init ${APP_NAME}; cd ${APP_NAME}
  INFO Using context "arn:aws:eks:eu-west-1:999999999999:cluster/training-gpu-1" from
 kubeconfig file "/home/ubuntu/.kube/config"
  ERROR Could not create app; directory '/home/ubuntu/kubeflow-tf-hvd' already exists
```

You can safely ignore this warning. However, you may have additional cleanup to do inside that folder. To simplify cleanup, delete the folder.

Usage Errors

```
ssh: Could not resolve hostname openmpi-worker-1.openmpi.kubeflow-dist-train-tf: Name or
 service not known
```

If you see this error message while using the Amazon EKS cluster, run the NVIDIA device plugin installation step again. Verify that you have targeted the right cluster by either passing in the specific config file or switching your active cluster to the targeted cluster.

Cleanup Errors

The following errors might be returned when cleaning up the resources of your Amazon EKS cluster.

- **Error: the server doesn't have a resource type "`namspace`"**

```
$ kubectl delete namespace ${NAMESPACE}
  error: the server doesn't have a resource type "namspace"
```

Verify the spelling of your namespace is correct.

- **Error: the server has asked for the client to provide credentials**

```
$ ks delete default
  ERROR the server has asked for the client to provide credentials
```

To solve this error, verify that `~/.kube/config` points to the correct cluster and that AWS credentials have been correctly configured using `aws configure` or by exporting AWS environment variables.

- **Error: finding app root from starting path: : unable to find ksonnet project**

```
$ ks delete default
  ERROR finding app root from starting path: : unable to find ksonnet project
```

To solve this error, verify that you are in the directory created by the ksonnet app. This is the folder where `ks init` was run.

- **Error: Error from server (NotFound): pods "openmpi-master" not found**

```
$ kubectl logs -n ${NAMESPACE} -f ${COMPONENT}-master > results/benchmark_1.out
  Error from server (NotFound): pods "openmpi-master" not found
```

This error might be caused by trying to access resources after the context is deleted. Deleting the default context causes the corresponding resources to be deleted as well.

Deep Learning Containers Images

AWS Deep Learning Containers are available as Docker images in Amazon ECR. Each Docker image is built for training or inference on a specific Deep Learning framework version, python version, with CPU or GPU support.

For the full list of available Deep Learning Containers and information on pulling them, see Available Deep Learning Containers Images.

Once you've selected your desired Deep Learning Containers image, continue with the one of the following:

- To run training and inference on Deep Learning Containers for Amazon EC2 using MXNet, PyTorch, TensorFlow, and TensorFlow 2, see Amazon EC2 Tutorials (p. 3)
- To run training and inference on Deep Learning Containers for Amazon ECS using MXNet, PyTorch, and TensorFlow, see Amazon ECS tutorials (p. 13)
- Deep Learning Containers for Amazon EKS offer CPU, GPU, and distributed GPU-based training, as well as CPU and GPU-based inference. To run training and inference on Deep Learning Containers for Amazon EKS using MXNet, PyTorch, and TensorFlow, see Amazon EKS Tutorials (p. 36)
- For information on security in Deep Learning Containers, see Security in AWS Deep Learning Containers (p. 90)
- For a list of the latest Deep Learning Containers release notes, see Release Notes for Deep Learning Containers (p. 97)

Deep Learning Containers Resources

The following topics describe additional AWS Deep Learning Containers resources.

Contents

-
-

Building AWS Deep Learning Containers Custom Images

How to Build Custom Images

We can easily customize both training and inference with Deep Learning Containers to add custom frameworks, libraries, and packages using Docker files.

Training with TensorFlow

In the following example Dockerfile, the resulting Docker image will have TensorFlow v1.15.2 optimized for GPUs and built to support Horovod and Python 3 for multi-node distributed training. It will also have the AWS samples GitHub repo which contains many deep learning model examples.

```
#Take base container
FROM 763104351884.dkr.ecr.us-east-1.amazonaws.com/tensorflow-training:1.15.2-gpu-py36-
cu100-ubuntu18.04

# Add custom stack of code
RUN git clone https://github.com/aws-samples/deep-learning-models
```

Training with Apache MXNet (Incubating)

In the following example Dockerfile, the resulting Docker image will have Apache MXNet (Incubating) v1.6.0 optimized for GPU inference built to support Horovod and Python 3. It will also have the MXNet GitHub repo which contains many deep learning model examples.

```
# Take the base MXNet Container
FROM 763104351884.dkr.ecr.us-east-1.amazonaws.com/mxnet-training:1.6.0-gpu-py36-cu101-
ubuntu16.04

# Add Custom stack of code
RUN git clone -b 1.6.0 https://github.com/apache/incubator-mxnet.git

ENTRYPOINT ["python", "/incubator-mxnet/example/image-classification/train_mnist.py"]
```

Build the Docker image, pointing to your personal Docker registry (usually your username), with the image's custom name and custom tag.

```
docker build -f Dockerfile -t <registry>/<any name>:<any tag>
```

Push to your personal Docker Registry:

```
docker push <registry>/<any name>:<any tag>
```

You can use the following command to run the container:

```
docker run -it < name or tag>
```

> **Important**
> You may need to login to access to the Deep Learning Containers image repository. Specify your region in the following command:

```
aws ecr get-login-password --region us-east-1 | docker login --username AWS --
password-stdin 763104351884.dkr.ecr.us-east-1.amazonaws.com
```

AWS Deep Learning Containers Intel Math Kernel Library (MKL) Recommendations

MKL Recommendation for CPU containers

Contents

The performance for training and inference workloads for a Deep Learning framework on CPU instances can vary and depend on a variety of configuration settings. As an example, on AWS EC2 c5.18xlarge instances the number of physical cores is 36 while the number of logical cores is 72. MKL's configuration settings for training and inference are influenced by these factors. By updating MKL's configuration to match your instance's capabilities, you may achieve performance improvements.

Consider the following examples using an Intel-MKL-optimized TensorFlow binary:

- A ResNet50v2 model, trained with TensorFlow and served for inference with TensorFlow Serving was observed to achieve 2x inference performance when the MKL settings were adjusted to match the instance's number cores. The following settings were used on a c5.18xlarge instance.

```
export TENSORFLOW_INTER_OP_PARALLELISM=2
# For an EC2 c5.18xlarge instance, number of logical cores = 72
export TENSORFLOW_INTRA_OP_PARALLELISM=72
# For an EC2 c5.18xlarge instance, number of physical cores = 36
export OMP_NUM_THREADS=36
export KMP_AFFINITY='granularity=fine,verbose,compact,1,0'
# For an EC2 c5.18xlarge instance, number of physical cores / 4 = 36 /4 = 9
```

```
export TENSORFLOW_SESSION_PARALLELISM=9
export KMP_BLOCKTIME=1
export KMP_SETTINGS=0
```

- A ResNet50_v1.5 model, trained with TensorFlow on the ImageNet dataset and using a NHWC image shape, the training throughput performance was observed to be around 9x faster. This is compared to the binary without MKL optimizations and measured in terms of samples/second. The following environment variables were used:

```
export TENSORFLOW_INTER_OP_PARALLELISM=0
# For an EC2 c5.18xlarge instance, number of logical cores = 72
export TENSORFLOW_INTRA_OP_PARALLELISM=72
# For an EC2 c5.18xlarge instance, number of physical cores = 36
export OMP_NUM_THREADS=36
export KMP_AFFINITY='granularity=fine,verbose,compact,1,0'
# For an EC2 c5.18xlarge instance, number of physical cores / 4 = 36 /4 = 9
export KMP_BLOCKTIME=1
export KMP_SETTINGS=0
```

The following links will help you learn how to use to tune Intel MKL and your Deep Learning framework's settings to optimize your deep learning workload:

- General best practices for Intel-optimized TensorFlow Serving

- TensorFlow performance

- Some Tips for improving Apache MXNet performance

- MXNet with Intel MKL-DNN - Performance Benchmarking

EC2 guide to set environment variables

Refer to `docker run` documentation on how to set environment variables when creating a container: https://docs.docker.com/engine/reference/run/#env-environment-variables

The following is an example on setting en environment variable called `OMP_NUM_THREADS` for docker run.

```
ubuntu@ip-172-31-95-248:~$ docker run -e OMP_NUM_THREADS=36 -it --entrypoint ""
 999999999999.dkr.ecr.us-east-1.amazonaws.com/beta-tensorflow-inference:1.13-py2-cpu-build
 bash
root@d437faf9b684:/# echo $OMP_NUM_THREADS
36
```

In rare cases Intel MKL can have adverse effects. To disable MKL with TensorFlow, set the following environment variables:

```
export TF_DISABLE_MKL=1
export TF_DISABLE_POOL_ALLOCATOR=1
```

ECS guide to set environment variables

To specify the environment variables for a container at runtime in ECS, you must edit the **ECS task definition**. Add the environment variables in the form of 'name' and 'value' key-pairs in `containerDefinitions` part of the task definition. The following is an example of setting `OMP_NUM_THREADS` and `KMP_BLOCKTIME` variables.

```
{
    "requiresCompatibilities": [
        "EC2"
    ],
    "containerDefinitions": [{
        "command": [
            "mkdir -p /test && cd /test && git clone -b r1.13 https://github.com/
tensorflow/serving.git && tensorflow_model_server --port=8500 --rest_api_port=8501
 --model_name=saved_model_half_plus_two_cpu --model_base_path=/test/serving/
tensorflow_serving/servables/tensorflow/testdata/saved_model_half_plus_two_cpu"
        ],
        "entryPoint": [
            "sh",
            "-c"
        ],
        "name": "EC2TFInference",
        "image": "999999999999.dkr.ecr.us-east-1.amazonaws.com/tf-inference:1.12-cpu-py3-
ubuntu16.04",
        "memory": 8111,
        "cpu": 256,
        "essential": true,
        "environment": [{
                "name": "OMP_NUM_THREADS",
                "value": "36"
            },
            {
                "name": "KMP_BLOCKTIME",
                "value": 1
            }
        ],
        "portMappings": [{
                "hostPort": 8500,
                "protocol": "tcp",
                "containerPort": 8500
            },
            {
                "hostPort": 8501,
                "protocol": "tcp",
                "containerPort": 8501
            },
            {
                "containerPort": 80,
                "protocol": "tcp"
            }
        ],
        "logConfiguration": {
            "logDriver": "awslogs",
            "options": {
                "awslogs-group": "/ecs/TFInference",
                "awslogs-region": "us-west-2",
                "awslogs-stream-prefix": "ecs",
                "awslogs-create-group": "true"
            }
        }
    }],
    "volumes": [],
```

```
        "networkMode": "bridge",
        "placementConstraints": [],
        "family": "Ec2TFInference"
}
```

In rare cases Intel MKL can have adverse effects. To disable MKL with TensorFlow, set the following environment variables:

```
        {
            "name": "TF_DISABLE_MKL",
            "value": 1
        },
        {
            "name": "TF_DISABLE_POOL_ALLOCATOR",
            "value": 1
        }
```

EKS guide to set environment variables

To specify the environment variables for the container at runtime, edit the raw manifests of your EKS job (.yaml, .json) . The following snippet of a manifest shows the definition of a container, with name `squeezenet-service`. Along with other attributes such as args and ports, the environment variables are listed in the form of 'name' and 'value' key-pairs.

```
      containers:
      - name: squeezenet-service
        image: 999999999999.dkr.ecr.us-east-1.amazonaws.com/beta-mxnet-inference:1.4.0-py3-
gpu-build
        command:
        - mxnet-model-server
        args:
        - --start
        - --mms-config /home/model-server/config.properties
        - --models squeezenet=https://s3.amazonaws.com/model-server/models/squeezenet_v1.1/
squeezenet_v1.1.model
        ports:
        - name: mms
          containerPort: 8080
        - name: mms-management
          containerPort: 8081
        imagePullPolicy: IfNotPresent
        env:
        - name: AWS_REGION
          value: us-east-1
        - name: OMP_NUM_THREADS
          value: 36
        - name: TENSORFLOW_INTER_OP_PARALLELISM
          value: 0
        - name: KMP_AFFINITY
          value: 'granularity=fine,verbose,compact,1,0'
        - name: KMP_BLOCKTIME
          value: 1
```

In rare cases Intel MKL can have adverse effects. To disable MKL with TensorFlow, set the following environment variables:

```
- name: TF_DISABLE_MKL
          value: 1
- name: TF_DISABLE_POOL_ALLOCATOR
```

```
            value: 1
```

Security in AWS Deep Learning Containers

Cloud security at AWS is the highest priority. As an AWS customer, you benefit from a data center and network architecture that is built to meet the requirements of the most security-sensitive organizations.

Security is a shared responsibility between AWS and you. The shared responsibility model describes this as security *of* the cloud and security *in* the cloud:

- **Security of the cloud** – AWS is responsible for protecting the infrastructure that runs AWS services in the AWS Cloud. AWS also provides you with services that you can use securely. Third-party auditors regularly test and verify the effectiveness of our security as part of the AWS Compliance Programs. To learn about the compliance programs that apply to Deep Learning Containers, see AWS Services in Scope by Compliance Program.
- **Security in the cloud** – Your responsibility is determined by the AWS service that you use. You are also responsible for other factors including the sensitivity of your data, your company's requirements, and applicable laws and regulations.

This documentation helps you understand how to apply the shared responsibility model when using Deep Learning Containers. The following topics show you how to configure Deep Learning Containers to meet your security and compliance objectives. You also learn how to use other AWS services that help you to monitor and secure your Deep Learning Containers resources.

For more information, see Security in Amazon EC2, Security in Amazon ECS, Security in Amazon EKS, and Security in Amazon SageMaker.

Topics

Data Protection in AWS Deep Learning Containers

The AWS shared responsibility model applies to data protection in AWS Deep Learning Containers. As described in this model, AWS is responsible for protecting the global infrastructure that runs all of the AWS Cloud. You are responsible for maintaining control over your content that is hosted on this infrastructure. This content includes the security configuration and management tasks for the AWS services that you use. For more information about data privacy, see the Data Privacy FAQ. For information about data protection in Europe, see the AWS Shared Responsibility Model and GDPR blog post on the *AWS Security Blog*.

For data protection purposes, we recommend that you protect AWS account credentials and set up individual user accounts with AWS Identity and Access Management (IAM). That way each user is given only the permissions necessary to fulfill their job duties. We also recommend that you secure your data in the following ways:

- Use multi-factor authentication (MFA) with each account.
- Use SSL/TLS to communicate with AWS resources. We recommend TLS 1.2 or later.
- Set up API and user activity logging with AWS CloudTrail.
- Use AWS encryption solutions, along with all default security controls within AWS services.
- Use advanced managed security services such as Amazon Macie, which assists in discovering and securing personal data that is stored in Amazon S3.
- If you require FIPS 140-2 validated cryptographic modules when accessing AWS through a command line interface or an API, use a FIPS endpoint. For more information about the available FIPS endpoints, see Federal Information Processing Standard (FIPS) 140-2.

We strongly recommend that you never put confidential or sensitive information, such as your customers' email addresses, into tags or free-form fields such as a **Name** field. This includes when you work with Deep Learning Containers or other AWS services using the console, API, AWS CLI, or AWS SDKs. Any data that you enter into tags or free-form fields used for names may be used for billing or diagnostic logs. If you provide a URL to an external server, we strongly recommend that you do not include credentials information in the URL to validate your request to that server.

Identity and Access Management in AWS Deep Learning Containers

AWS Identity and Access Management (IAM) is an AWS service that helps an administrator securely control access to AWS resources. IAM administrators control who can be *authenticated* (signed in) and *authorized* (have permissions) to use Deep Learning Containers resources. IAM is an AWS service that you can use with no additional charge.

For more information on Identity and Access Management, see Identity and Access Management for Amazon EC2, Identity and Access Management for Amazon ECS, Identity and Access Management for Amazon EKS, and Identity and Access Management for Amazon SageMaker.

Topics

- Authenticating With Identities (p. 91)
- Managing Access Using Policies (p. 93)
- IAM with Amazon EMR (p. 95)

Authenticating With Identities

Authentication is how you sign in to AWS using your identity credentials. For more information about signing in using the AWS Management Console, see Signing in to the AWS Management Console as an IAM user or root user in the *IAM User Guide*.

You must be *authenticated* (signed in to AWS) as the AWS account root user, an IAM user, or by assuming an IAM role. You can also use your company's single sign-on authentication or even sign in using Google or Facebook. In these cases, your administrator previously set up identity federation using IAM roles. When you access AWS using credentials from another company, you are assuming a role indirectly.

To sign in directly to the AWS Management Console, use your password with your root user email address or your IAM user name. You can access AWS programmatically using your root user or IAM users access keys. AWS provides SDK and command line tools to cryptographically sign your request using your credentials. If you don't use AWS tools, you must sign the request yourself. Do this using

Signature Version 4, a protocol for authenticating inbound API requests. For more information about authenticating requests, see Signature Version 4 signing process in the *AWS General Reference*.

Regardless of the authentication method that you use, you might also be required to provide additional security information. For example, AWS recommends that you use multi-factor authentication (MFA) to increase the security of your account. To learn more, see Using multi-factor authentication (MFA) in AWS in the *IAM User Guide*.

AWS account root user

When you first create an AWS account, you begin with a single sign-in identity that has complete access to all AWS services and resources in the account. This identity is called the AWS account *root user* and is accessed by signing in with the email address and password that you used to create the account. We strongly recommend that you do not use the root user for your everyday tasks, even the administrative ones. Instead, adhere to the best practice of using the root user only to create your first IAM user. Then securely lock away the root user credentials and use them to perform only a few account and service management tasks.

IAM Users and Groups

An *IAM user* is an identity within your AWS account that has specific permissions for a single person or application. An IAM user can have long-term credentials such as a user name and password or a set of access keys. To learn how to generate access keys, see Managing access keys for IAM users in the *IAM User Guide*. When you generate access keys for an IAM user, make sure you view and securely save the key pair. You cannot recover the secret access key in the future. Instead, you must generate a new access key pair.

An *IAM group* is an identity that specifies a collection of IAM users. You can't sign in as a group. You can use groups to specify permissions for multiple users at a time. Groups make permissions easier to manage for large sets of users. For example, you could have a group named *IAMAdmins* and give that group permissions to administer IAM resources.

Users are different from roles. A user is uniquely associated with one person or application, but a role is intended to be assumable by anyone who needs it. Users have permanent long-term credentials, but roles provide temporary credentials. To learn more, see When to create an IAM user (instead of a role) in the *IAM User Guide*.

IAM Roles

An *IAM role* is an identity within your AWS account that has specific permissions. It is similar to an IAM user, but is not associated with a specific person. You can temporarily assume an IAM role in the AWS Management Console by switching roles. You can assume a role by calling an AWS CLI or AWS API operation or by using a custom URL. For more information about methods for using roles, see Using IAM roles in the *IAM User Guide*.

IAM roles with temporary credentials are useful in the following situations:

- **Temporary IAM user permissions** – An IAM user can assume an IAM role to temporarily take on different permissions for a specific task.
- **Federated user access** – Instead of creating an IAM user, you can use existing identities from AWS Directory Service, your enterprise user directory, or a web identity provider. These are known as *federated users*. AWS assigns a role to a federated user when access is requested through an identity provider. For more information about federated users, see Federated users and roles in the *IAM User Guide*.
- **Cross-account access** – You can use an IAM role to allow someone (a trusted principal) in a different account to access resources in your account. Roles are the primary way to grant cross-account access. However, with some AWS services, you can attach a policy directly to a resource (instead of using a role

as a proxy). To learn the difference between roles and resource-based policies for cross-account access, see How IAM roles differ from resource-based policies in the *IAM User Guide*.

- **Cross-service access** – Some AWS services use features in other AWS services. For example, when you make a call in a service, it's common for that service to run applications in Amazon EC2 or store objects in Amazon S3. A service might do this using the calling principal's permissions, using a service role, or using a service-linked role.

  - **Principal permissions** – When you use an IAM user or role to perform actions in AWS, you are considered a principal. Policies grant permissions to a principal. When you use some services, you might perform an action that then triggers another action in a different service. In this case, you must have permissions to perform both actions. To see whether an action requires additional dependent actions in a policy, see in the *Service Authorization Reference*.

  - **Service role** – A service role is an IAM role that a service assumes to perform actions on your behalf. An IAM administrator can create, modify, and delete a service role from within IAM. For more information, see Creating a role to delegate permissions to an AWS service in the *IAM User Guide*.

  - **Service-linked role** – A service-linked role is a type of service role that is linked to an AWS service. The service can assume the role to perform an action on your behalf. Service-linked roles appear in your IAM account and are owned by the service. An IAM administrator can view, but not edit the permissions for service-linked roles.

- **Applications running on Amazon EC2** – You can use an IAM role to manage temporary credentials for applications that are running on an EC2 instance and making AWS CLI or AWS API requests. This is preferable to storing access keys within the EC2 instance. To assign an AWS role to an EC2 instance and make it available to all of its applications, you create an instance profile that is attached to the instance. An instance profile contains the role and enables programs that are running on the EC2 instance to get temporary credentials. For more information, see Using an IAM role to grant permissions to applications running on Amazon EC2 instances in the *IAM User Guide*.

To learn whether to use IAM roles or IAM users, see When to create an IAM role (instead of a user) in the *IAM User Guide*.

Managing Access Using Policies

You control access in AWS by creating policies and attaching them to IAM identities or AWS resources. A policy is an object in AWS that, when associated with an identity or resource, defines their permissions. You can sign in as the root user or an IAM user, or you can assume an IAM role. When you then make a request, AWS evaluates the related identity-based or resource-based policies. Permissions in the policies determine whether the request is allowed or denied. Most policies are stored in AWS as JSON documents. For more information about the structure and contents of JSON policy documents, see Overview of JSON policies in the *IAM User Guide*.

Administrators can use AWS JSON policies to specify who has access to what. That is, which **principal** can perform **actions** on what **resources**, and under what **conditions**.

Every IAM entity (user or role) starts with no permissions. In other words, by default, users can do nothing, not even change their own password. To give a user permission to do something, an administrator must attach a permissions policy to a user. Or the administrator can add the user to a group that has the intended permissions. When an administrator gives permissions to a group, all users in that group are granted those permissions.

IAM policies define permissions for an action regardless of the method that you use to perform the operation. For example, suppose that you have a policy that allows the `iam:GetRole` action. A user with that policy can get role information from the AWS Management Console, the AWS CLI, or the AWS API.

Identity-Based Policies

Identity-based policies are JSON permissions policy documents that you can attach to an identity, such as an IAM user, group of users, or role. These policies control what actions users and roles can perform,

on which resources, and under what conditions. To learn how to create an identity-based policy, see Creating IAM policies in the *IAM User Guide*.

Identity-based policies can be further categorized as *inline policies* or *managed policies*. Inline policies are embedded directly into a single user, group, or role. Managed policies are standalone policies that you can attach to multiple users, groups, and roles in your AWS account. Managed policies include AWS managed policies and customer managed policies. To learn how to choose between a managed policy or an inline policy, see Choosing between managed policies and inline policies in the *IAM User Guide*.

Resource-Based Policies

Resource-based policies are JSON policy documents that you attach to a resource. Examples of resource-based policies are IAM *role trust policies* and Amazon S3 *bucket policies*. In services that support resource-based policies, service administrators can use them to control access to a specific resource. For the resource where the policy is attached, the policy defines what actions a specified principal can perform on that resource and under what conditions. You must specify a principal in a resource-based policy. Principals can include accounts, users, roles, federated users, or AWS services.

Resource-based policies are inline policies that are located in that service. You can't use AWS managed policies from IAM in a resource-based policy.

Access Control Lists (ACLs)

Access control lists (ACLs) control which principals (account members, users, or roles) have permissions to access a resource. ACLs are similar to resource-based policies, although they do not use the JSON policy document format.

Amazon S3, AWS WAF, and Amazon VPC are examples of services that support ACLs. To learn more about ACLs, see Access control list (ACL) overview in the *Amazon Simple Storage Service Developer Guide*.

Other Policy Types

AWS supports additional, less-common policy types. These policy types can set the maximum permissions granted to you by the more common policy types.

- **Permissions boundaries** – A permissions boundary is an advanced feature in which you set the maximum permissions that an identity-based policy can grant to an IAM entity (IAM user or role). You can set a permissions boundary for an entity. The resulting permissions are the intersection of entity's identity-based policies and its permissions boundaries. Resource-based policies that specify the user or role in the `Principal` field are not limited by the permissions boundary. An explicit deny in any of these policies overrides the allow. For more information about permissions boundaries, see Permissions boundaries for IAM entities in the *IAM User Guide*.

- **Service control policies (SCPs)** – SCPs are JSON policies that specify the maximum permissions for an organization or organizational unit (OU) in AWS Organizations. AWS Organizations is a service for grouping and centrally managing multiple AWS accounts that your business owns. If you enable all features in an organization, then you can apply service control policies (SCPs) to any or all of your accounts. The SCP limits permissions for entities in member accounts, including each AWS account root user. For more information about Organizations and SCPs, see How SCPs work in the *AWS Organizations User Guide*.

- **Session policies** – Session policies are advanced policies that you pass as a parameter when you programmatically create a temporary session for a role or federated user. The resulting session's permissions are the intersection of the user or role's identity-based policies and the session policies. Permissions can also come from a resource-based policy. An explicit deny in any of these policies overrides the allow. For more information, see Session policies in the *IAM User Guide*.

Multiple Policy Types

When multiple types of policies apply to a request, the resulting permissions are more complicated to understand. To learn how AWS determines whether to allow a request when multiple policy types are involved, see Policy evaluation logic in the *IAM User Guide*.

IAM with Amazon EMR

You can use AWS Identity and Access Management with Amazon EMR to define users, AWS resources, groups, roles, and policies. You can also control which AWS services these users and roles can access.

For more information on using IAM with Amazon EMR, see AWS Identity and Access Management for Amazon EMR.

Logging and Monitoring in AWS Deep Learning Containers

Your AWS Deep Learning Containers does not come with monitoring utilities. For information on monitoring, see GPU Monitoring and Optimization, Monitoring Amazon EC2, Monitoring Amazon ECS, Monitoring Amazon EKS, and Monitoring Amazon SageMaker.

Usage Tracking

The following Deep Learning Containers include code that allows AWS to collect the instance types, frameworks, framework versions, container types, and Python versions used for the containers. No information on the commands used within the containers is collected or retained. No other information about the containers is collected or retained.

- TensorFlow 1.15
- TensorFlow 2.0
- TensorFlow 2.1
- PyTorch 1.2
- PyTorch 1.3.1
- MXNet 1.6

To opt-out of usage tracking, use a custom entrypoint to disable the call for the following services:

- Amazon EC2 Custom Entrypoints
- Amazon ECS Custom Entrypoints
- Amazon EKS Custom Entrypoints

To opt out of usage tracking for TensorFlow (version >=1.15), PyTorch (version >= 1.5), and MXNet (version >= 1.6) containers, you also need to set the OPT_OUT_TRACKING environment variable.

```
OPT_OUT_TRACKING=true
```

Compliance Validation for AWS Deep Learning Containers

Third-party auditors assess the security and compliance of services as part of multiple AWS compliance programs. For information on the supported compliance programs, see Compliance Validation for Amazon EC2, Compliance Validation for Amazon ECS, Compliance Validation for Amazon EKS, and Compliance Validation for Amazon SageMaker.

For a list of AWS services in scope of specific compliance programs, see AWS Services in Scope by Compliance Program. For general information, see AWS Compliance Programs.

You can download third-party audit reports using AWS Artifact. For more information, see Downloading Reports in AWS Artifact.

Your compliance responsibility when using Deep Learning Containers is determined by the sensitivity of your data, your company's compliance objectives, and applicable laws and regulations. AWS provides the following resources to help with compliance:

- Security and Compliance Quick Start Guides – These deployment guides discuss architectural considerations and provide steps for deploying security- and compliance-focused baseline environments on AWS.
- AWS Compliance Resources – This collection of workbooks and guides might apply to your industry and location.
- Evaluating Resources with Rules in the *AWS Config Developer Guide* – The AWS Config service assesses how well your resource configurations comply with internal practices, industry guidelines, and regulations.
- AWS Security Hub – This AWS service provides a comprehensive view of your security state within AWS that helps you check your compliance with security industry standards and best practices.

Resilience in AWS Deep Learning Containers

The AWS global infrastructure is built around AWS Regions and Availability Zones. AWS Regions provide multiple physically separated and isolated Availability Zones, which are connected with low-latency, high-throughput, and highly redundant networking. With Availability Zones, you can design and operate applications and databases that automatically fail over between zones without interruption. Availability Zones are more highly available, fault tolerant, and scalable than traditional single or multiple data center infrastructures.

For more information about AWS Regions and Availability Zones, see AWS Global Infrastructure.

For information on features to help support your data resiliency and backup needs, see Resilience in Amazon EC2, Resilience in Amazon EKS, and Resilience in Amazon SageMaker.

Infrastructure Security in AWS Deep Learning Containers

The infrastructure security of AWS Deep Learning Containers is backed by Amazon EC2, Amazon ECS, Amazon EKS, or SageMaker. For more information, see Infrastructure Security in Amazon EC2, Infrastructure Security in Amazon ECS, Infrastructure Security in Amazon EKS, and Infrastructure Security in Amazon SageMaker.

Release Notes for Deep Learning Containers

Check the latest release notes for AWS Deep Learning Containers built for specific machine learning frameworks, infrastructures, and AWS services.

Note
Starting with MXNet 1.9, PyTorch 1.10, and Tensorflow 2.7, CPU and GPU Deep Learning Containers are released as either **E3** or **SageMaker** images. E3 images are supported on Amazon EC2, Amazon ECS, and Amazon EKS, whereas SageMaker images are supported on Amazon SageMaker and contain additional packages to support SageMaker workloads.

Single-framework Deep Learning Containers

TensorFlow

- AWS Deep Learning Containers for TensorFlow 2.9.0 (E3): May 19, 2022
- AWS Deep Learning Containers for TensorFlow 2.8.0 (E3): March 22, 2022
- AWS Deep Learning Containers for TensorFlow 2.8.0 (SageMaker): March 22, 2022
- AWS Deep Learning Containers for TensorFlow 2.7.0 (E3): December 15, 2021
- AWS Deep Learning Containers for TensorFlow 2.7.0 (SageMaker): March 22, 2022
- AWS Deep Learning Containers for TensorFlow 2.6.0: September 24, 2021
- AWS Deep Learning Containers for TensorFlow 2.5.0: July 01, 2021
- AWS Deep Learning Containers for TensorFlow 2.4.1: March 15, 2021
- AWS Deep Learning Containers for TensorFlow 2.3.2: March 15, 2021
- AWS Deep Learning Containers for TensorFlow 2.2.2: March 15, 2021
- AWS Deep Learning Containers for TensorFlow 2.1.3: March 15, 2021
- AWS Deep Learning Containers for TensorFlow 2.0.4: March 15, 2021
- AWS Deep Learning Containers for TensorFlow 2.3.1 with CUDA 11.0: October 15, 2020
- AWS Deep Learning Containers with Elastic Inference for TensorFlow 2.0.0: August 31, 2020
- AWS Deep Learning Containers with Elastic Inference for TensorFlow 1.15.0: August 31, 2020
- AWS Deep Learning Containers for TensorFlow 2.3.0: August 7, 2020
- AWS Deep Learning Containers for TensorFlow 2.2.0 with CUDA 10.2: July 24, 2020
- AWS Deep Learning Containers for TensorFlow 2.2.0 with CUDA 10.1: July 20, 2020
- AWS Deep Learning Containers for Tensorflow 1.15.3 July 29, 2020
- AWS Deep Learning Containers for TensorFlow 2.1.1: June 25, 2020
- AWS Deep Learning Containers for TensorFlow 2.0.2: June 19, 2020
- AWS Deep Learning Containers for TensorFlow 2.2.0 with CUDA 10.1: May 20, 2020
- AWS Deep Learning Containers for Tensorflow 2.1.0: March 19, 2020
- AWS Deep Learning Containers for TensorFlow 2.0: February 26, 2020
- AWS Deep Learning Containers for Tensorflow 1.15 with python-3.7 support: May 6, 2020
- AWS Deep Learning Containers for Tensorflow 1.15.2 March 19, 2020

PyTorch

- AWS Deep Learning Containers for PyTorch 1.11.0 (SageMaker): May 06, 2022
- AWS Deep Learning Containers for PyTorch 1.11.0 (E3): April 14, 2022
- AWS Deep Learning Containers for PyTorch 1.10.0 (E3): November 3, 2021
- AWS Deep Learning Containers for PyTorch 1.10.2 (SageMaker): April 14, 2022
- AWS Deep Learning Containers for PyTorch 1.9.0: August 30, 2021
- AWS Deep Learning Containers for PyTorch 1.8.0: March 16, 2021
- AWS Deep Learning Containers for PyTorch 1.7.1 with CUDA 11.0: March 15, 2021
- AWS Deep Learning Containers for PyTorch 1.6.0 with CUDA 11.0: December 9, 2020
- AWS Deep Learning Containers with Elastic Inference for PyTorch 1.3.1: August 31, 2020
- AWS Deep Learning Containers for PyTorch 1.6.0: August 3, 2020
- AWS Deep Learning Containers for PyTorch 1.5.1: June 19, 2020
- AWS Deep Learning Containers for PyTorch 1.4: June 06, 2020
- AWS Deep Learning Containers for PyTorch 1.5.0: May 05, 2020
- AWS Deep Learning Containers for PyTorch 1.4: April 03, 2020

AWS MX powered by Apache MXNet

- AWS Deep Learning Containers for MXNet 1.9.0 (E3): January 31, 2022
- AWS Deep Learning Containers for MXNet 1.9.0 (SageMaker): January 31, 2022
- AWS Deep Learning Containers for MXNet 1.8.0: December 01, 2020
- AWS Deep Learning Containers for MXNet 1.7.0: September 17, 2020
- AWS Deep Learning Containers with Elastic Inference for MXNet 1.5.1: August 31, 2020
- AWS Deep Learning Containers for MXNet 1.6.0: July 20, 2020
- AWS Deep Learning Containers for MXNet 1.6.0: April 03, 2020

Graviton Deep Learning Containers

TensorFlow

- AWS Deep Learning Containers for Graviton TensorFlow 2.7.0 (E3): December 04, 2021

PyTorch

- AWS Deep Learning Containers for Graviton PyTorch 1.10.0 (E3): December 04, 2021

Habana Deep Learning Containers

TensorFlow

- AWS Deep Learning Containers for Habana TensorFlow 2.5.0: October 26, 2021

PyTorch

- AWS Deep Learning Containers for Habana PyTorch 1.7.1: October 26, 2021

Document History for Deep Learning Containers Developer Guide

The following table describes the documentation for this release of Deep Learning Containers.

- **API version: latest**
- **Latest documentation update:** February 26, 2020

update-history-change	update-history-description	update-history-date
Apache MXNet with Horovod (p. 1)	Apache MXNet tutorial was added to the developer guide.	February 26, 2020
Deep Learning Containers Developer Guide Launch (p. 1)	Deep Learning Containers setup and tutorials were added to the developer guide.	February 17, 2020

AWS glossary

For the latest AWS terminology, see the AWS glossary in the *AWS General Reference*.